

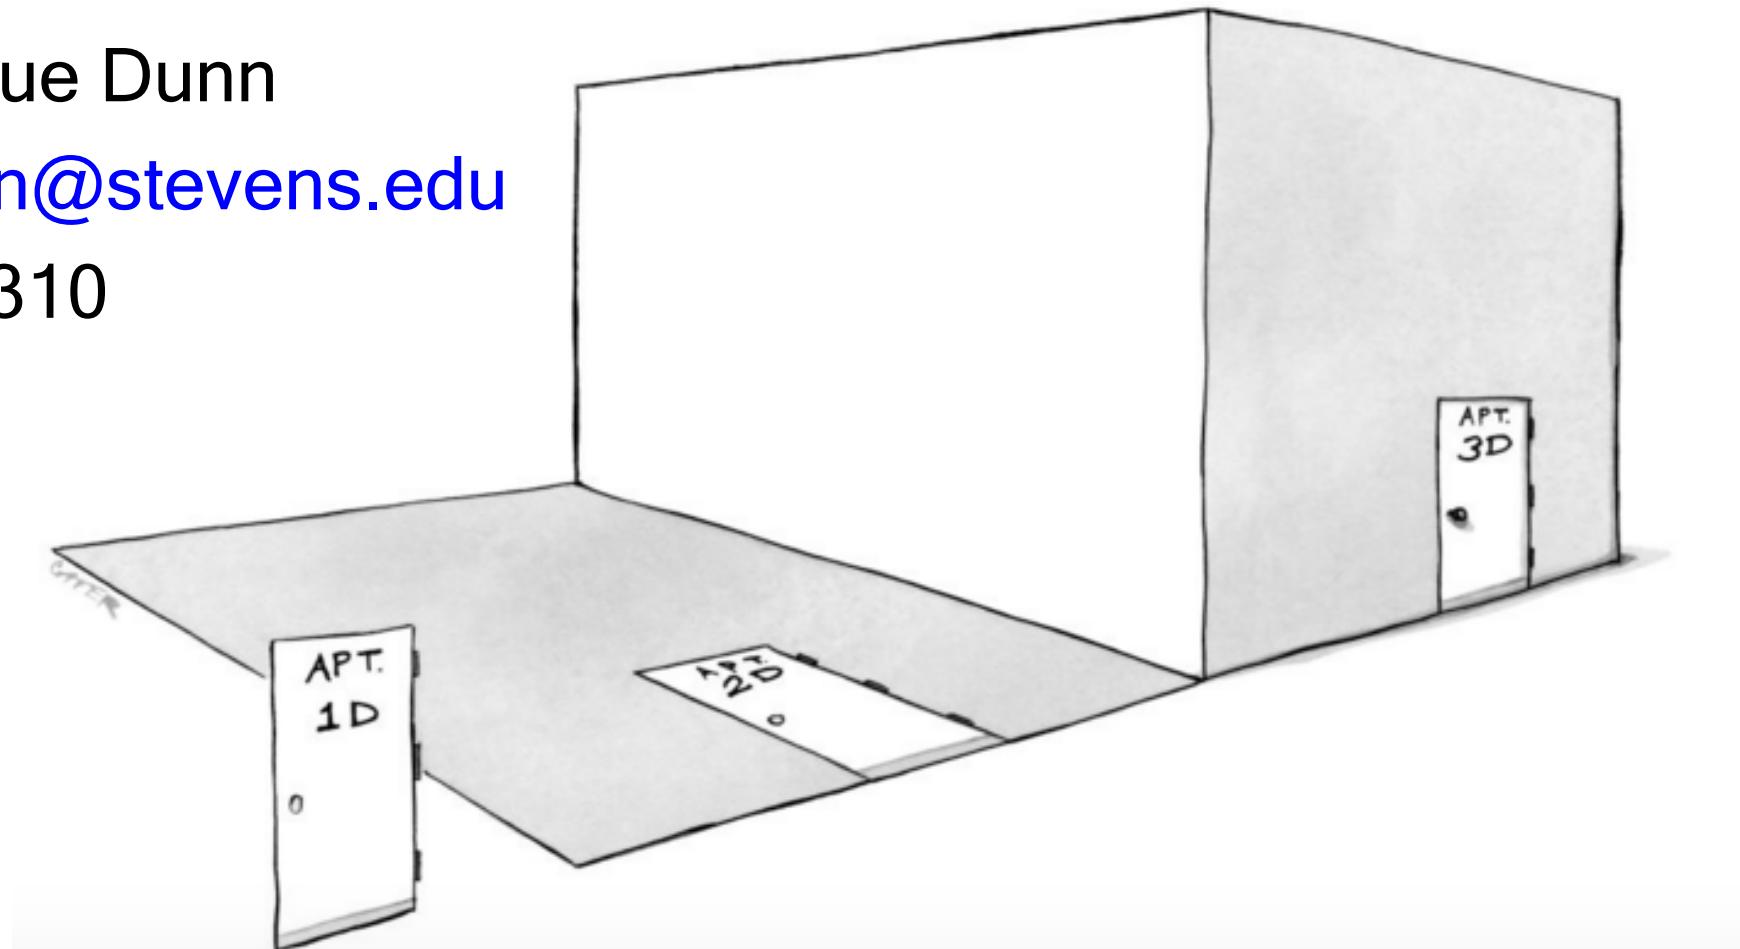
# CS 532: 3D Computer Vision

## Lecture 11

Enrique Dunn

[edunn@stevens.edu](mailto:edunn@stevens.edu)

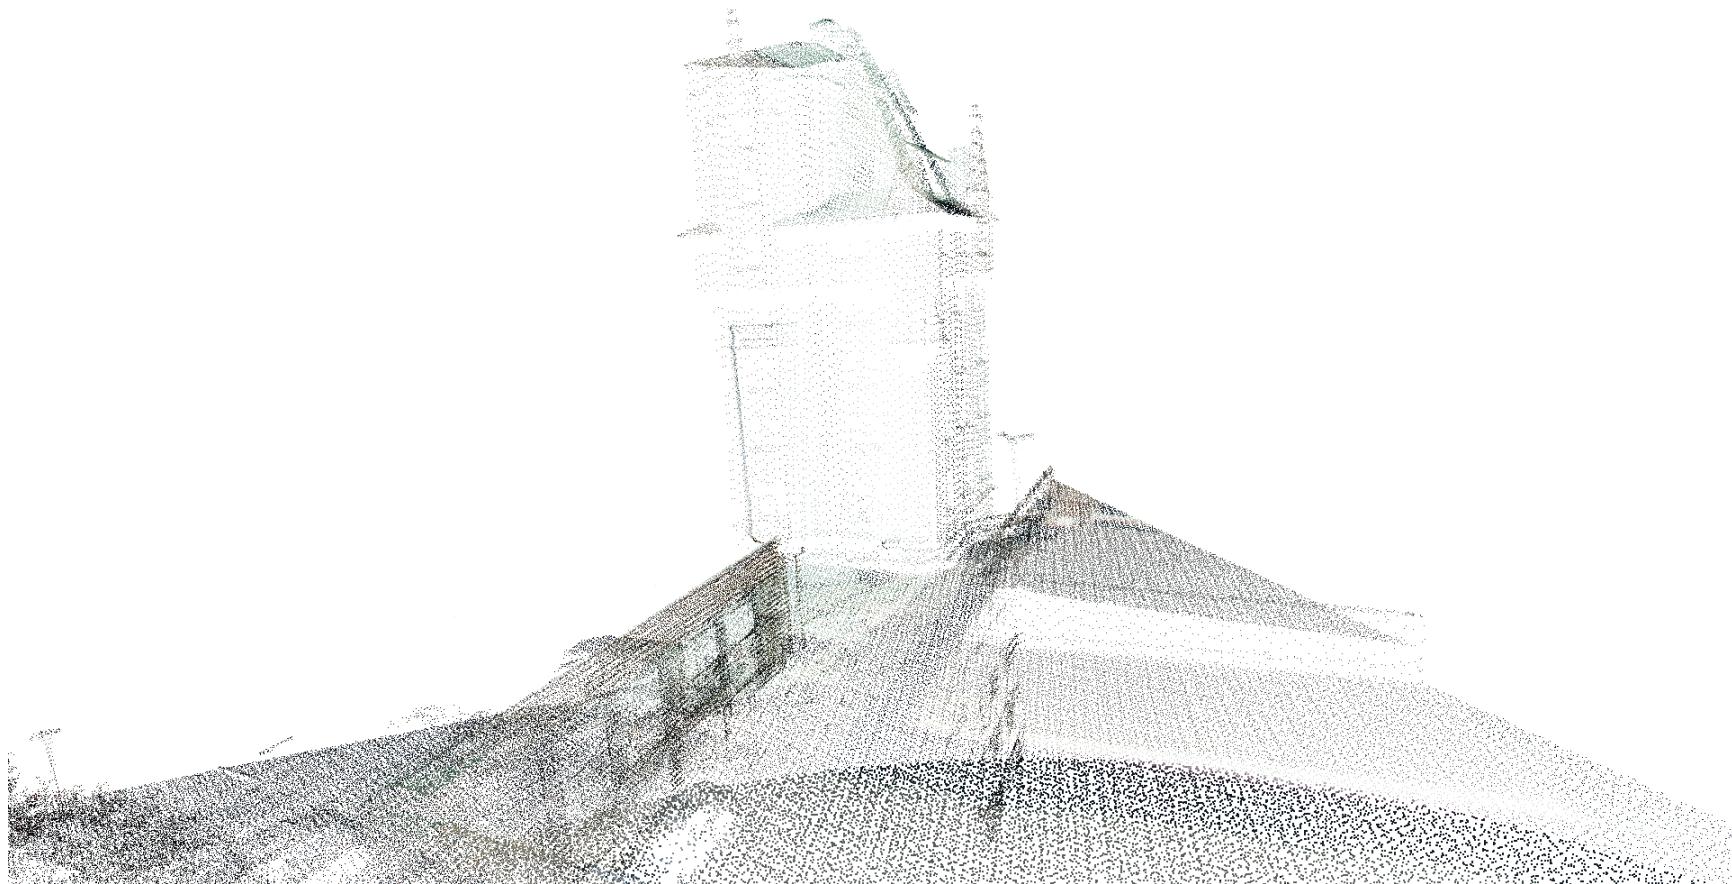
Lieb 310



# Lecture Outline

- Unorganized point clouds
  - Range queries and kd-trees
- Descriptors for 3D point clouds
- David M. Mount, CMSC 754:  
Computational Geometry lecture notes,  
Department of Computer Science,  
University of Maryland, Spring 2012
  - Lecture 16

# Point Clouds



# Unorganized Point Clouds

- In many cases, data come as unorganized point clouds
  - Not trivial to generate meshes even given a 2D reference frame (see Homework 6.2)
  - Even more complicated to fuse multiple viewpoint-based meshes
  - Some sensors (moving LIDAR) measure only points

# What can be done?

- While mesh generation is possible, it is not guaranteed to be correct
  - Even watertight manifolds can be generated using Poisson surface reconstruction (see previous notes)
  - It is often hard to discriminate forward and backward facing points
  - It is hard to determine whether there are gaps between points

# What can be done?

- Process the data by assuming possible connections between points
  - I.e. assuming nearby points belong to the same smooth surface
- Do not make hard decisions

# Key Parameter: Scale

- Trade-off between fidelity to data and robustness to noise
- If analysis is performed at small scale, details (high curvature) can be preserved
- If data are noisy, high curvature typically corresponds to noise...

# Nearest Neighbors

- Number one enabling technology for working with unorganized point clouds: **Nearest Neighbor Search**
- Algorithms cannot be quadratic
- It can be shown that kd-trees can be generated in  $O(n \log n)$  time
- Queries take  $O(\log n)$  per point
  - k queries: find k nearest neighbors of query point
  - epsilon queries: find neighbors within an  $\epsilon$ -ball centered at the query point

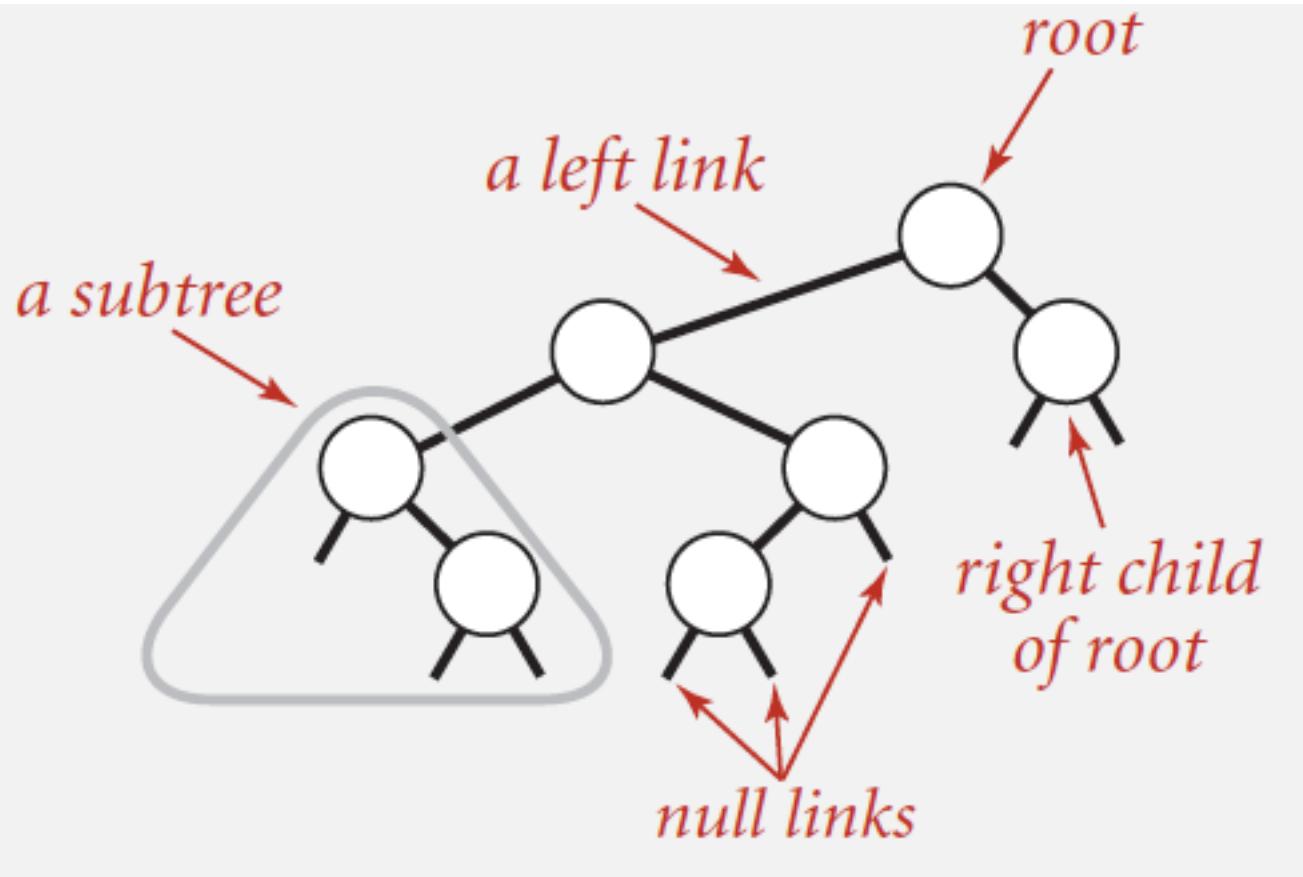
# Tools

- knnsearch() in Matlab
  - <http://www.mathworks.com/help/stats/knnsearch.html>
  - Do multiple queries with one call
- FLANN library
  - <http://www.cs.ubc.ca/research/flann/>
  - Bindings for C/C++, Matlab and python, part of PCL
- CvKNearest in OpenCV (which has bindings for C/C++, Python and Java)
- KDTree in scikit
- All these are not limited to 3D data

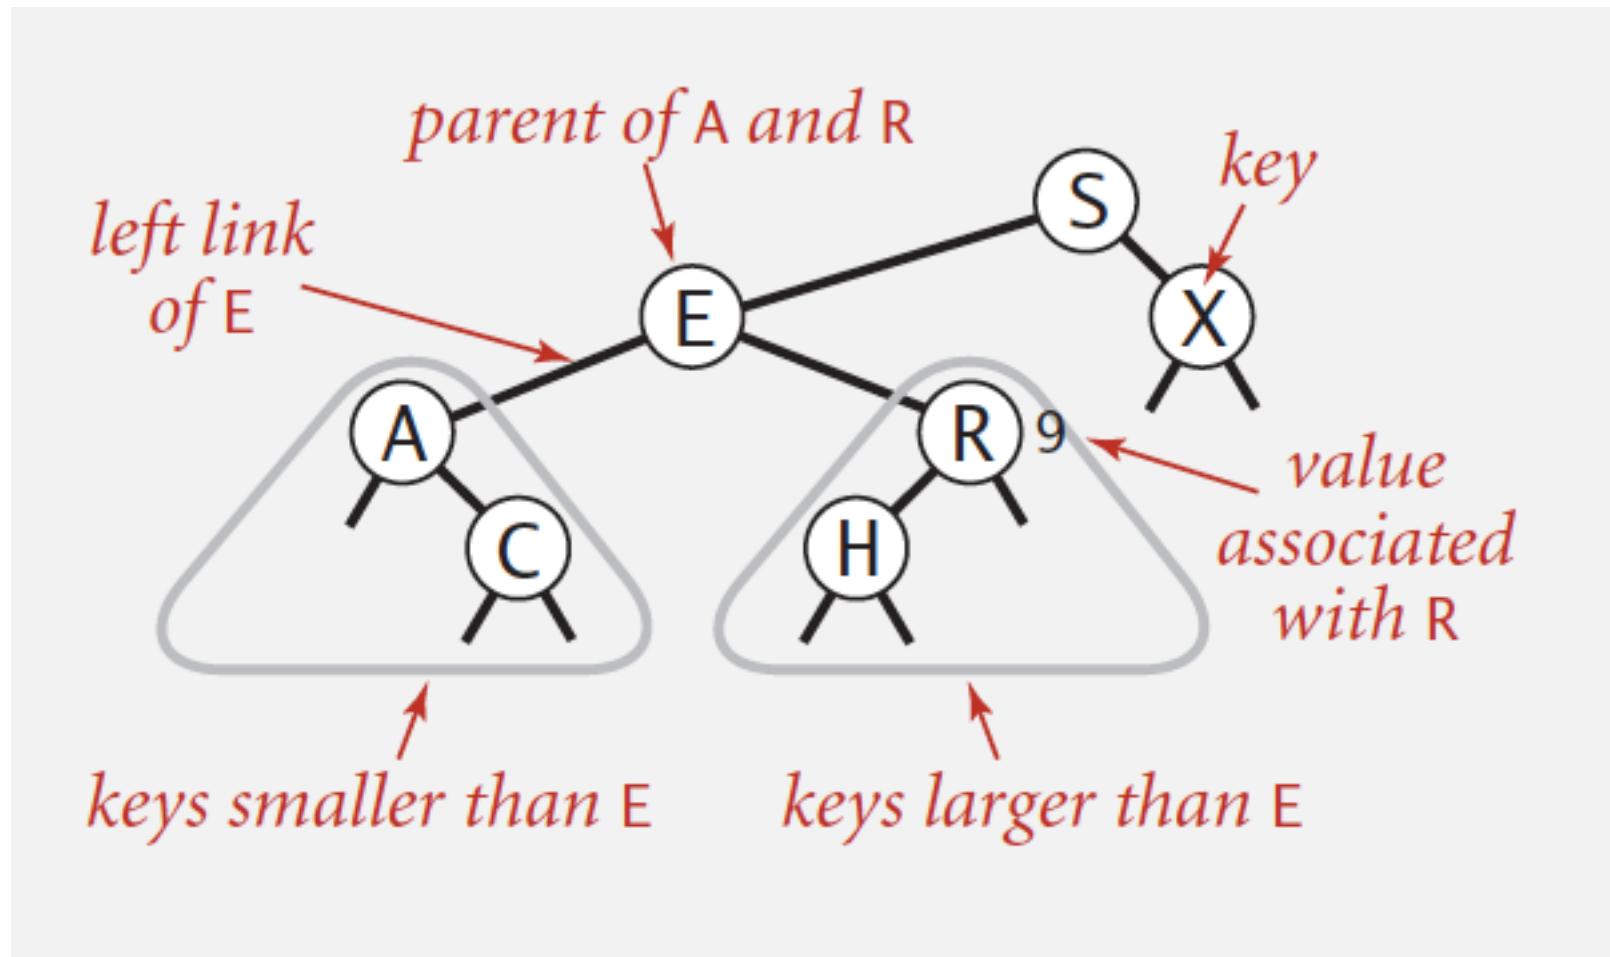
# 1D Range Queries

- 1D range query problem: Preprocess a set of  $n$  points on the real line such that the ones inside a 1D query range (interval) can be reported fast
- The points  $p_1, \dots, p_n$  are known beforehand, the query  $[x, x']$  only later
- A solution to a query problem is a data structure description, a query algorithm, and a construction algorithm

# Binary Search Tree

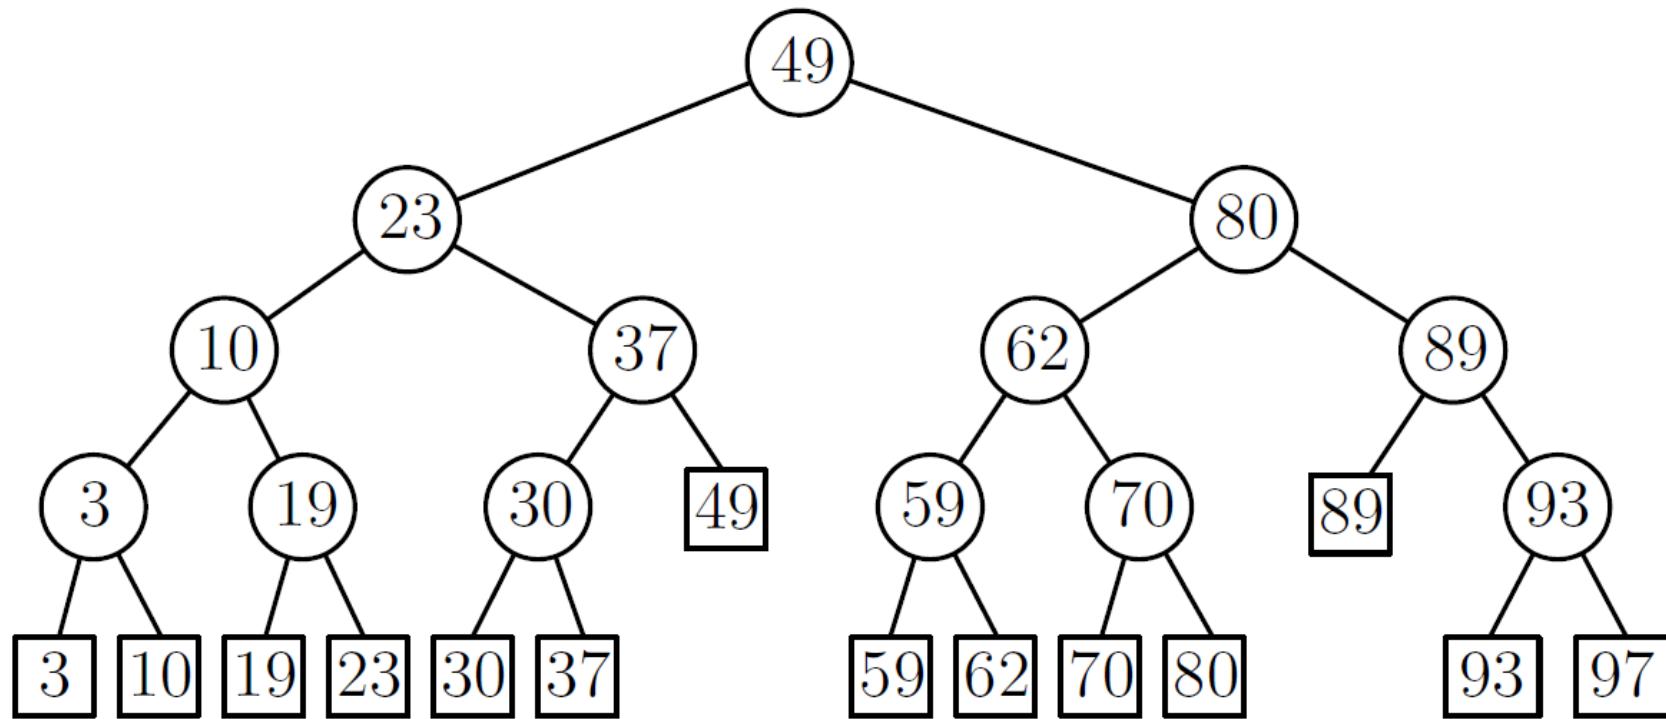


# Binary Search Tree



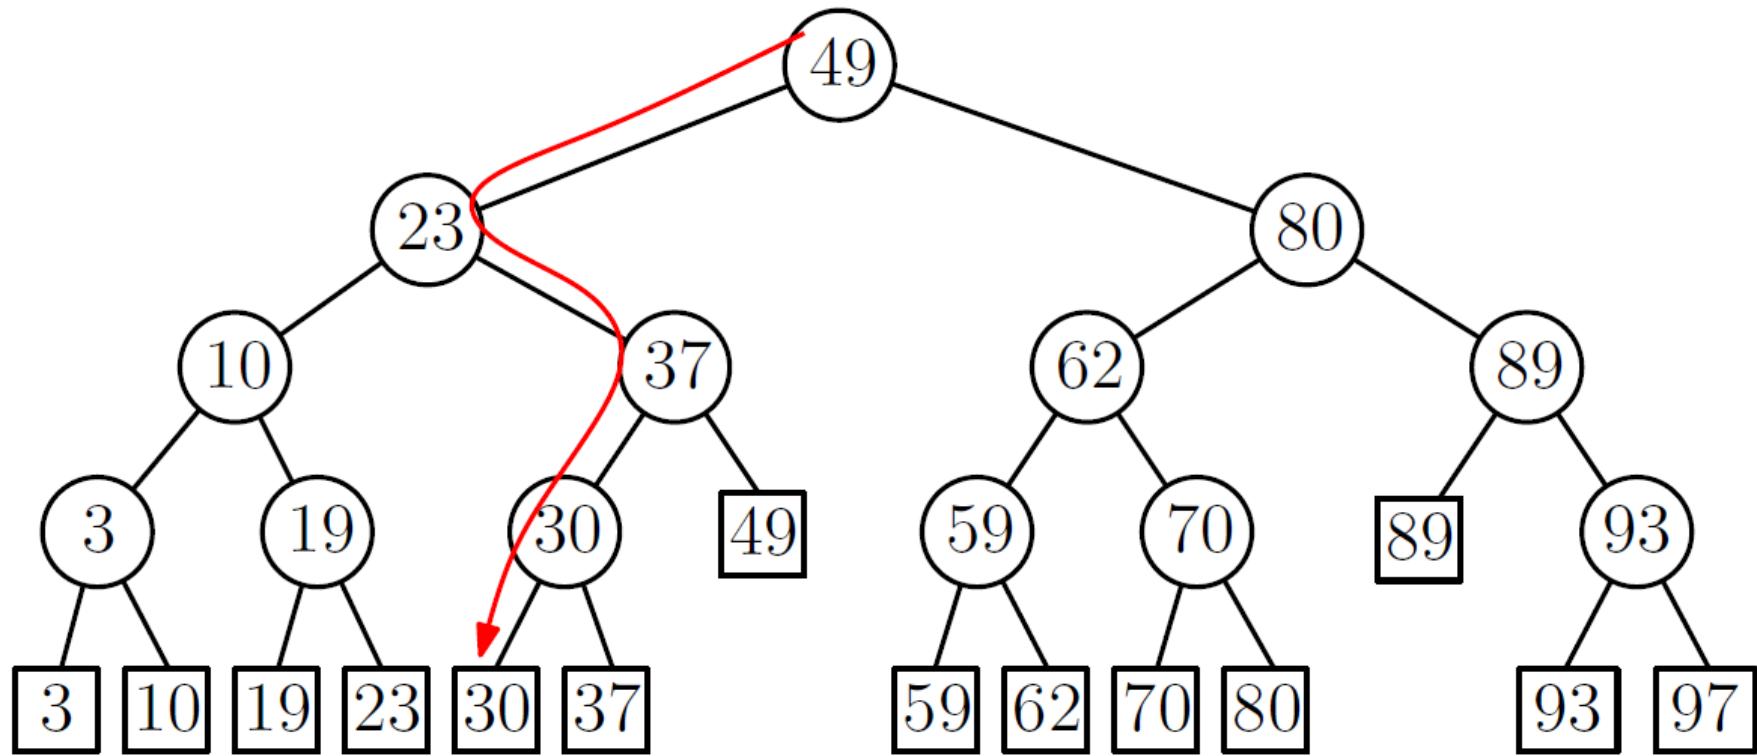
# Balanced Binary Search Trees

- A balanced binary search tree with the points in the leaves



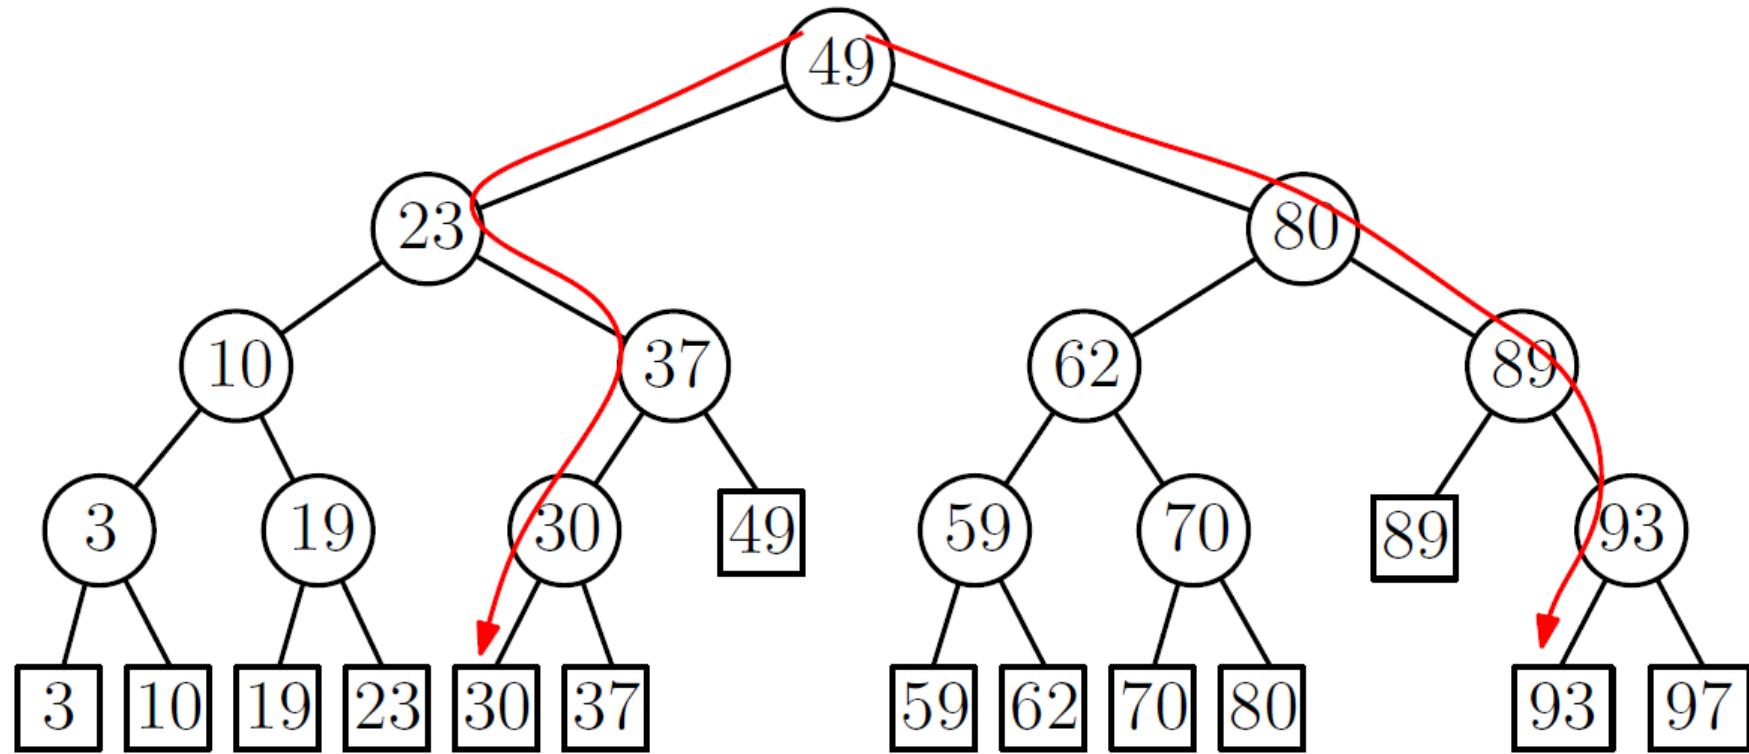
# Balanced Binary Search Trees

- The search path for 25



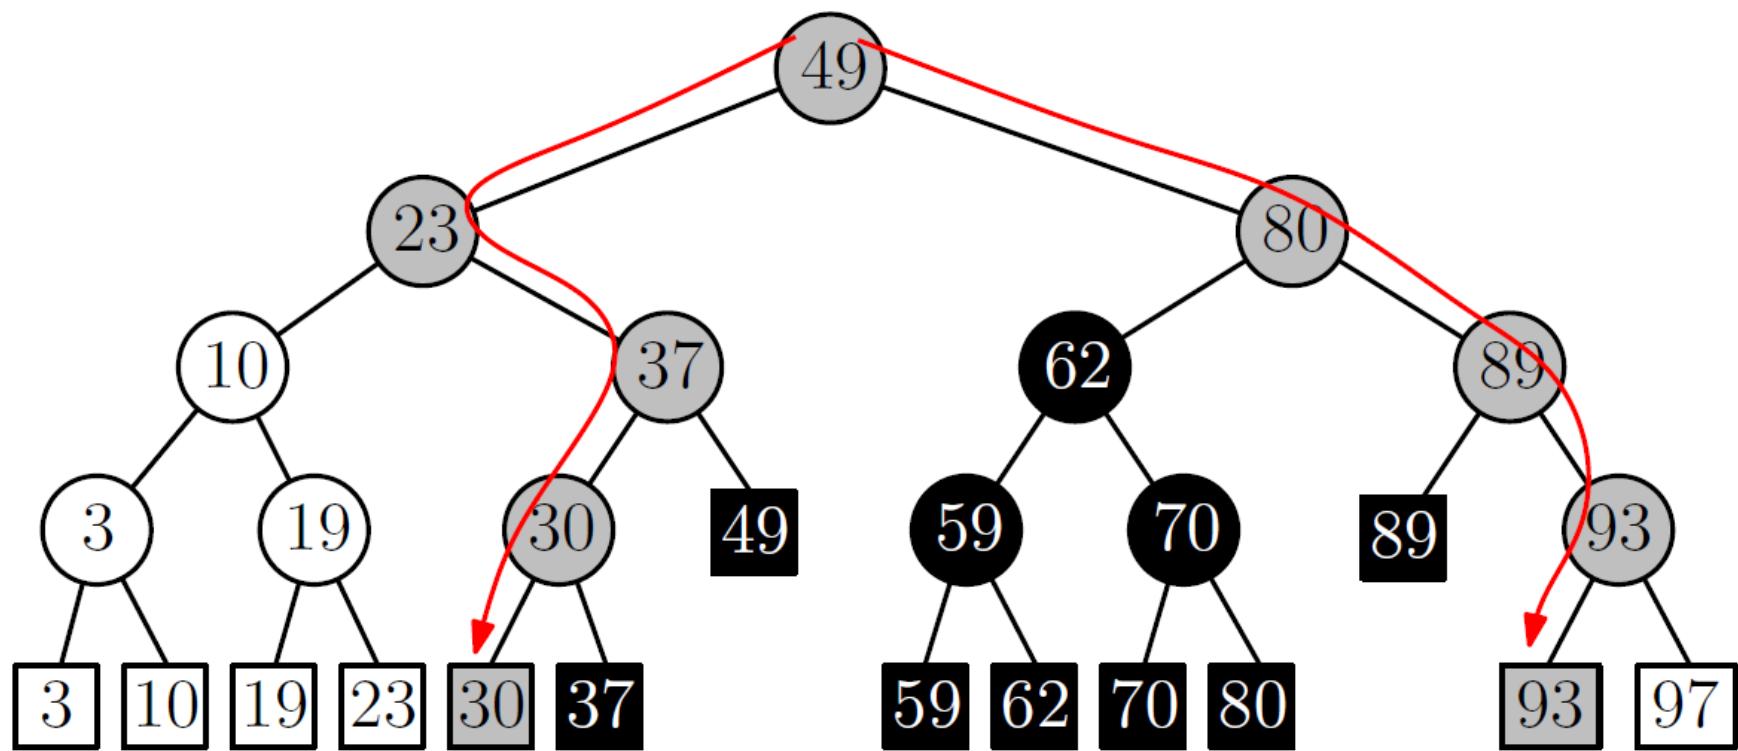
# Balanced Binary Search Trees

- The search paths for 25 and 90



# Example 1D Range Query

- A 1D range query with [25, 90]



# Node Types for a Query

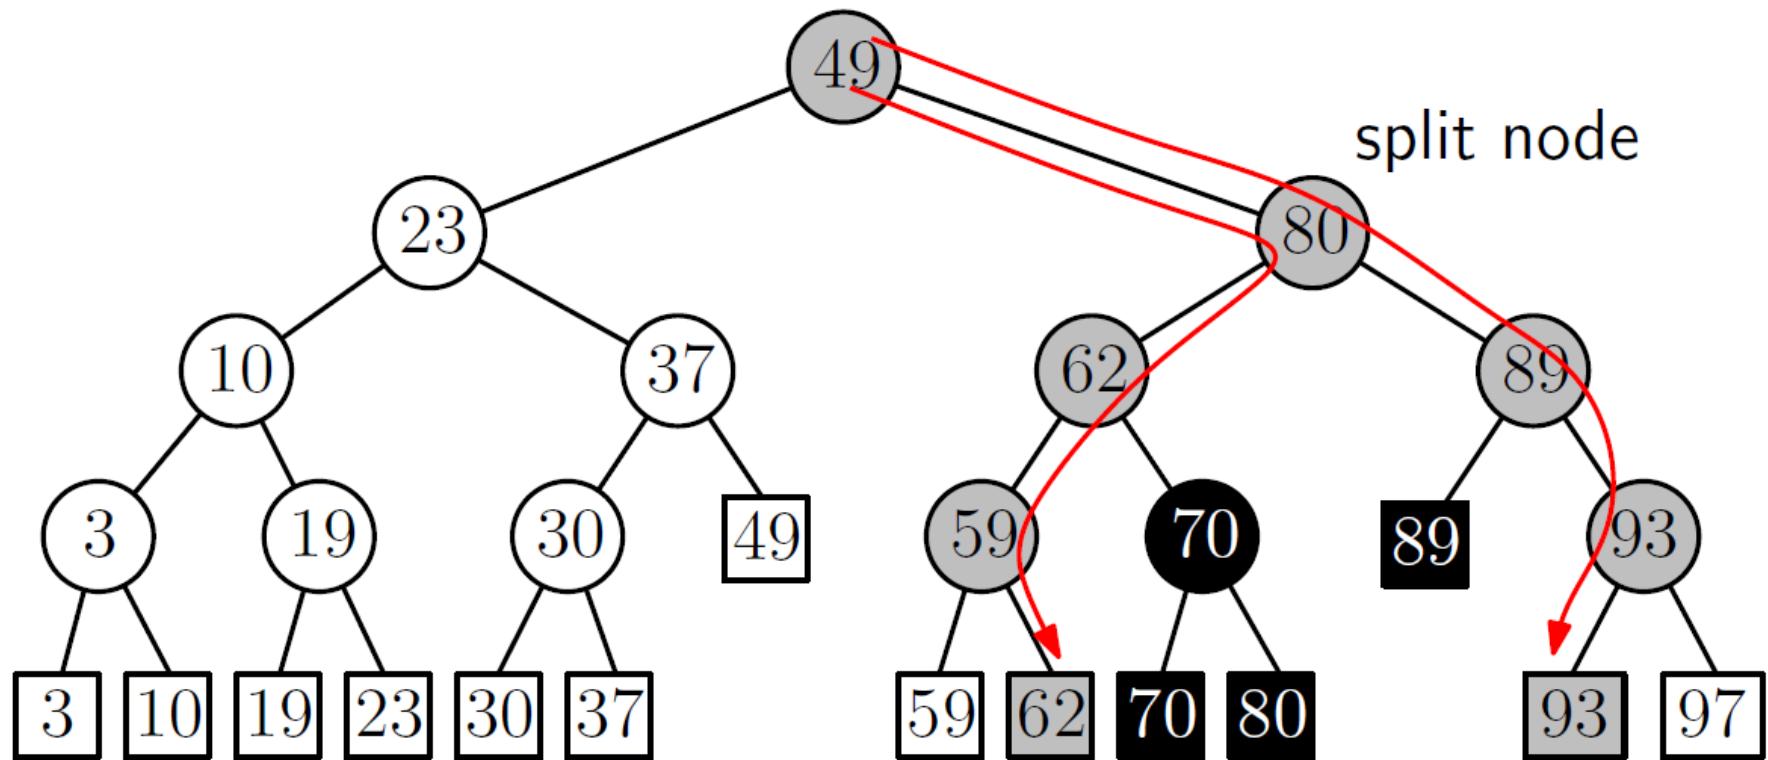
- Three types of nodes for a given query:
  - White nodes: never visited by the query
  - Grey nodes: visited by the query, unclear if they lead to output
  - Black nodes: visited by the query, whole subtree is output

# Node Types for a Query

- The query algorithm comes down to what we do at each type of node
- Grey nodes: use query range to decide how to proceed: to not visit a subtree (pruning), to report a complete subtree, or just continue
- Black nodes: traverse and enumerate all points in the leaves

# Example 1D Range Query

- A 1D range query with [61, 90]



# 1D Range Query Algorithm

**Algorithm** 1DRANGEQUERY( $\mathcal{T}, [x : x']$ )

1.  $v_{\text{split}} \leftarrow \text{FINDSPLITNODE}(\mathcal{T}, x, x')$
2. **if**  $v_{\text{split}}$  is a leaf
3.     **then** Check if the point in  $v_{\text{split}}$  must be reported.
4.     **else**  $v \leftarrow lc(v_{\text{split}})$
5.         **while**  $v$  is not a leaf
6.             **do if**  $x \leq x_v$
7.                 **then** REPORTSUBTREE( $rc(v)$ )
8.                  $v \leftarrow lc(v)$
9.                 **else**  $v \leftarrow rc(v)$
10.         Check if the point stored in  $v$  must be reported.
11.          $v \leftarrow rc(v_{\text{split}})$
12.         Similarly, follow the path to  $x'$ , and ...

# Query Time Analysis

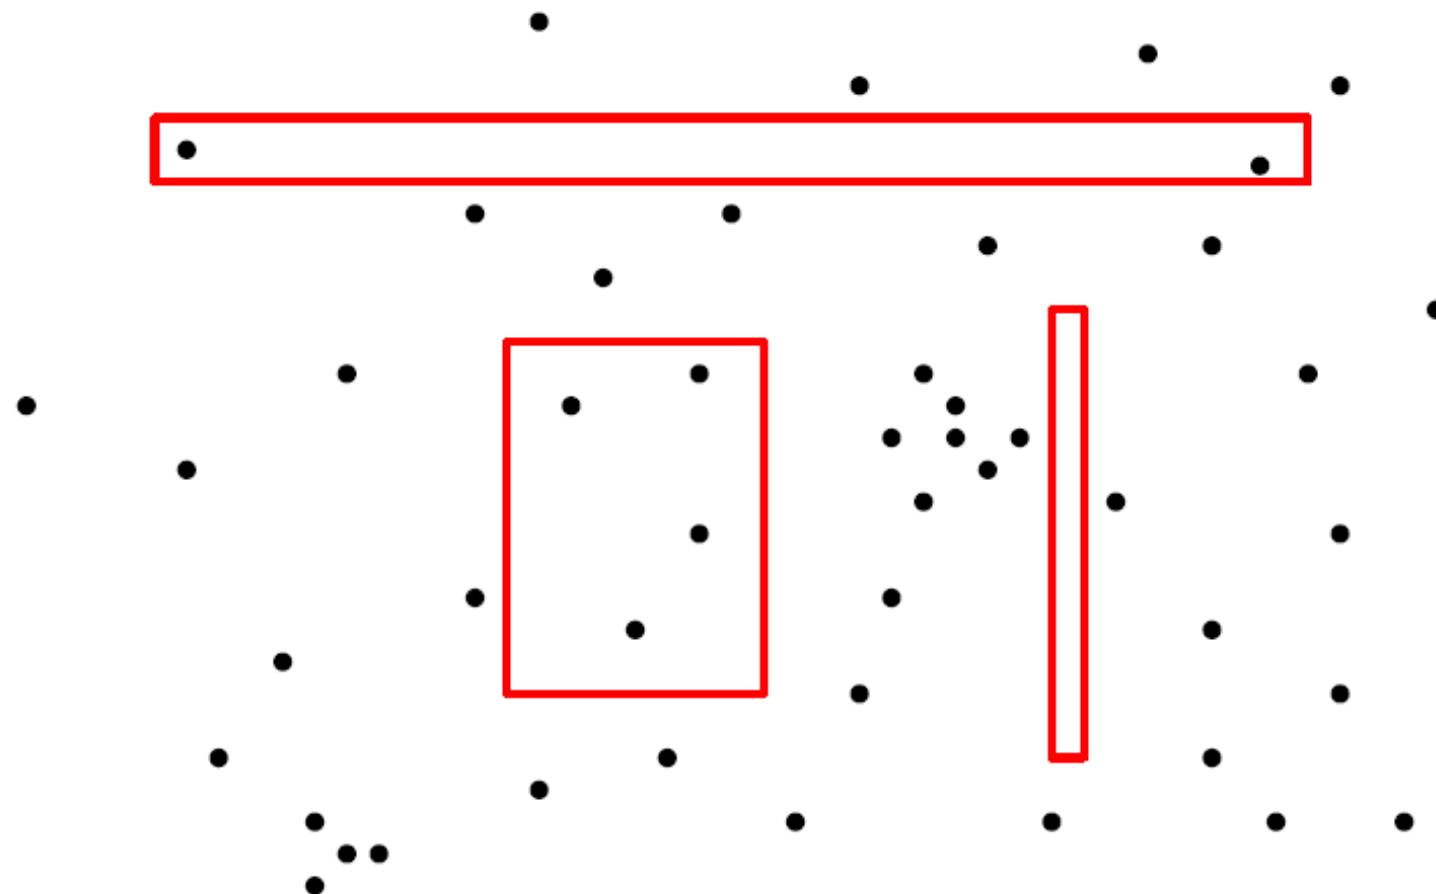
- The efficiency analysis is based on counting the numbers of nodes visited for each type
  - White nodes: never visited by the query; no time spent
  - Grey nodes: visited by the query, unclear if they lead to output; time complexity depends on  $n$
  - Black nodes: visited by the query, whole subtree is output; time complexity depends on  $k$ , the output size

# Query Time Analysis

- Grey nodes: they occur on only two paths in the tree, and since the tree is balanced, its depth is  $O(\log n)$
- Black nodes: a (sub)tree with  $m$  leaves has  $m-1$  internal nodes; traversal visits  $O(m)$  nodes and finds  $m$  points for the output
- The time spent at each node is  $O(1)$

=> $O(\log n + k)$  query time

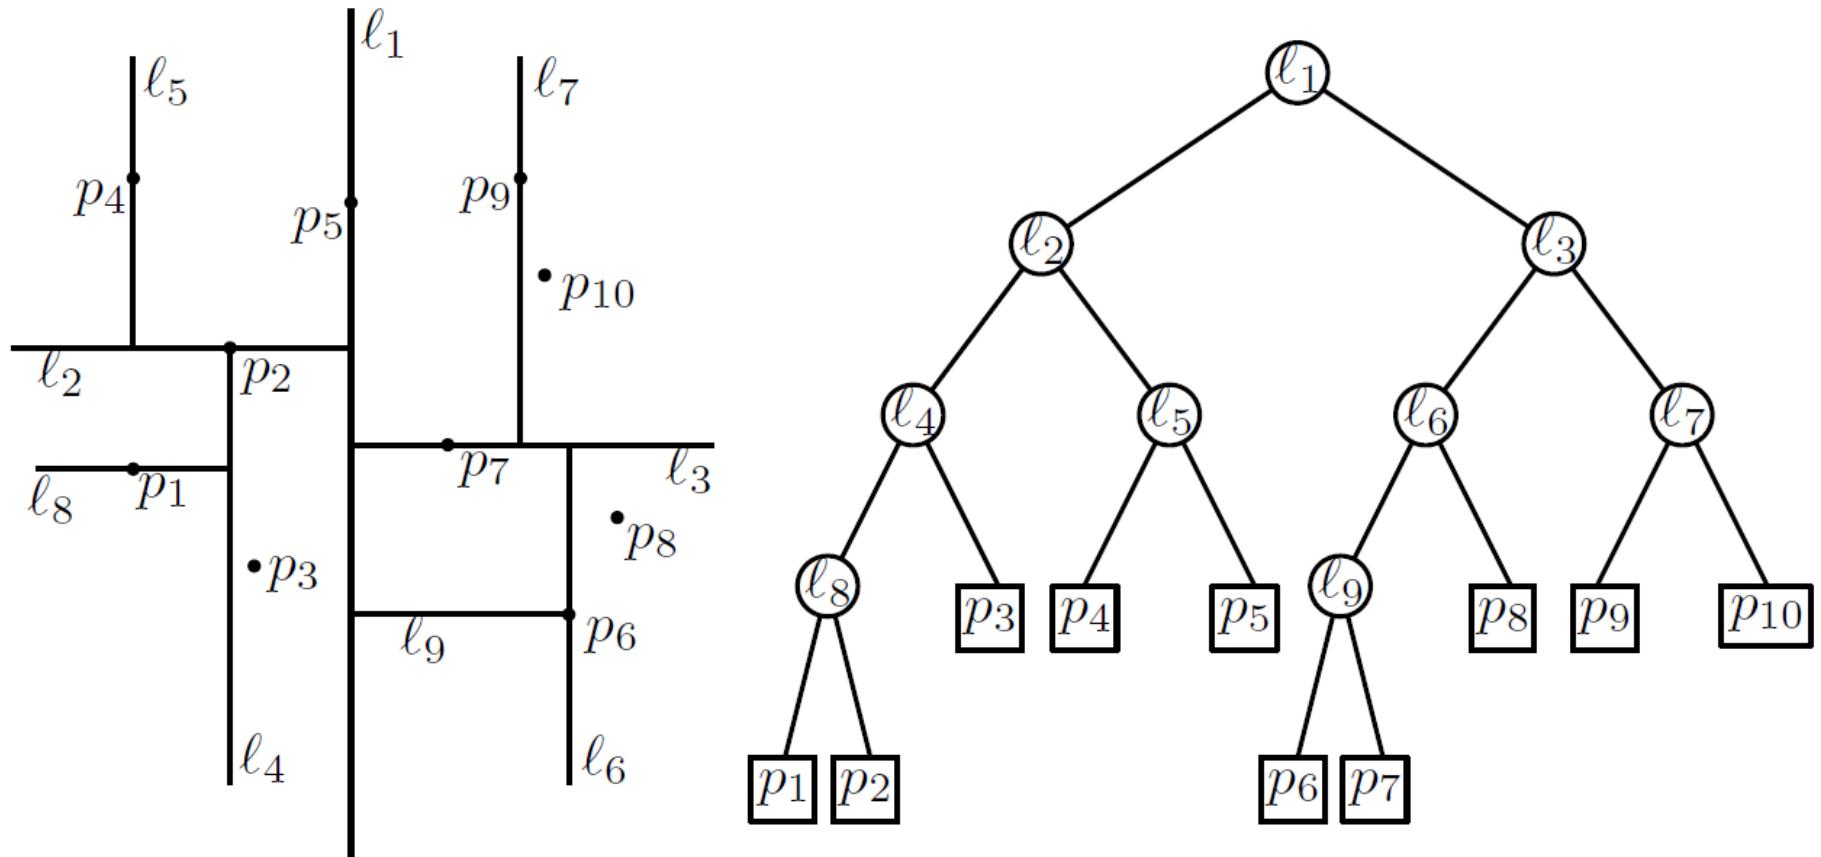
# Range Queries in 2D



# Kd-Trees

- Kd-trees, the idea: Split the point set alternately by x-coordinate and by y-coordinate
- Split by x-coordinate: split by a vertical line that has half the points to its left or on it, and half to its right
- Split by y-coordinate: split by a horizontal line that has half the points below or on it, and half above it

# Kd-Trees

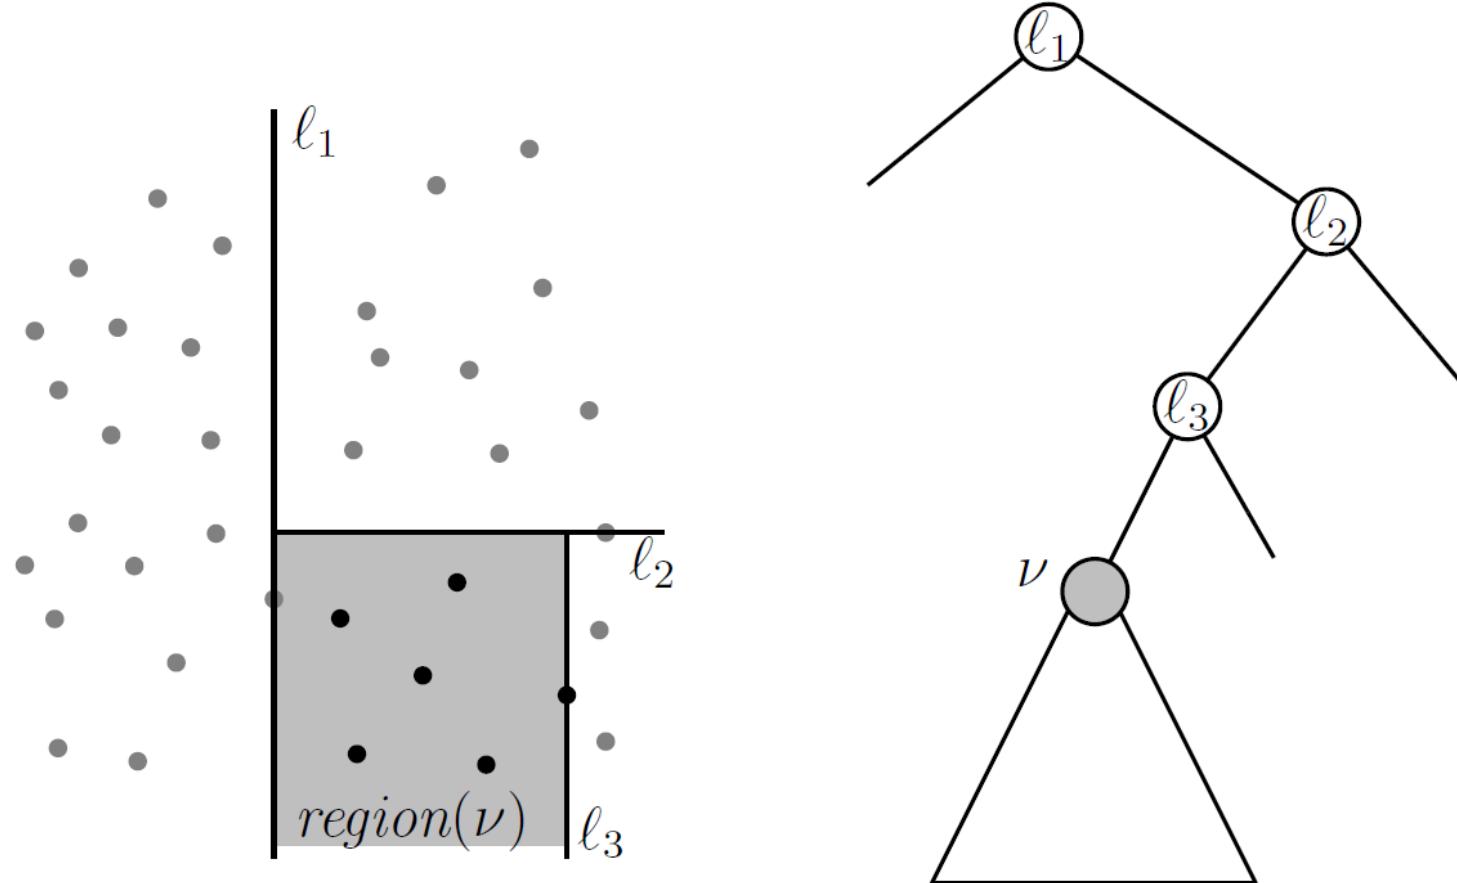


# Kd-Tree Construction

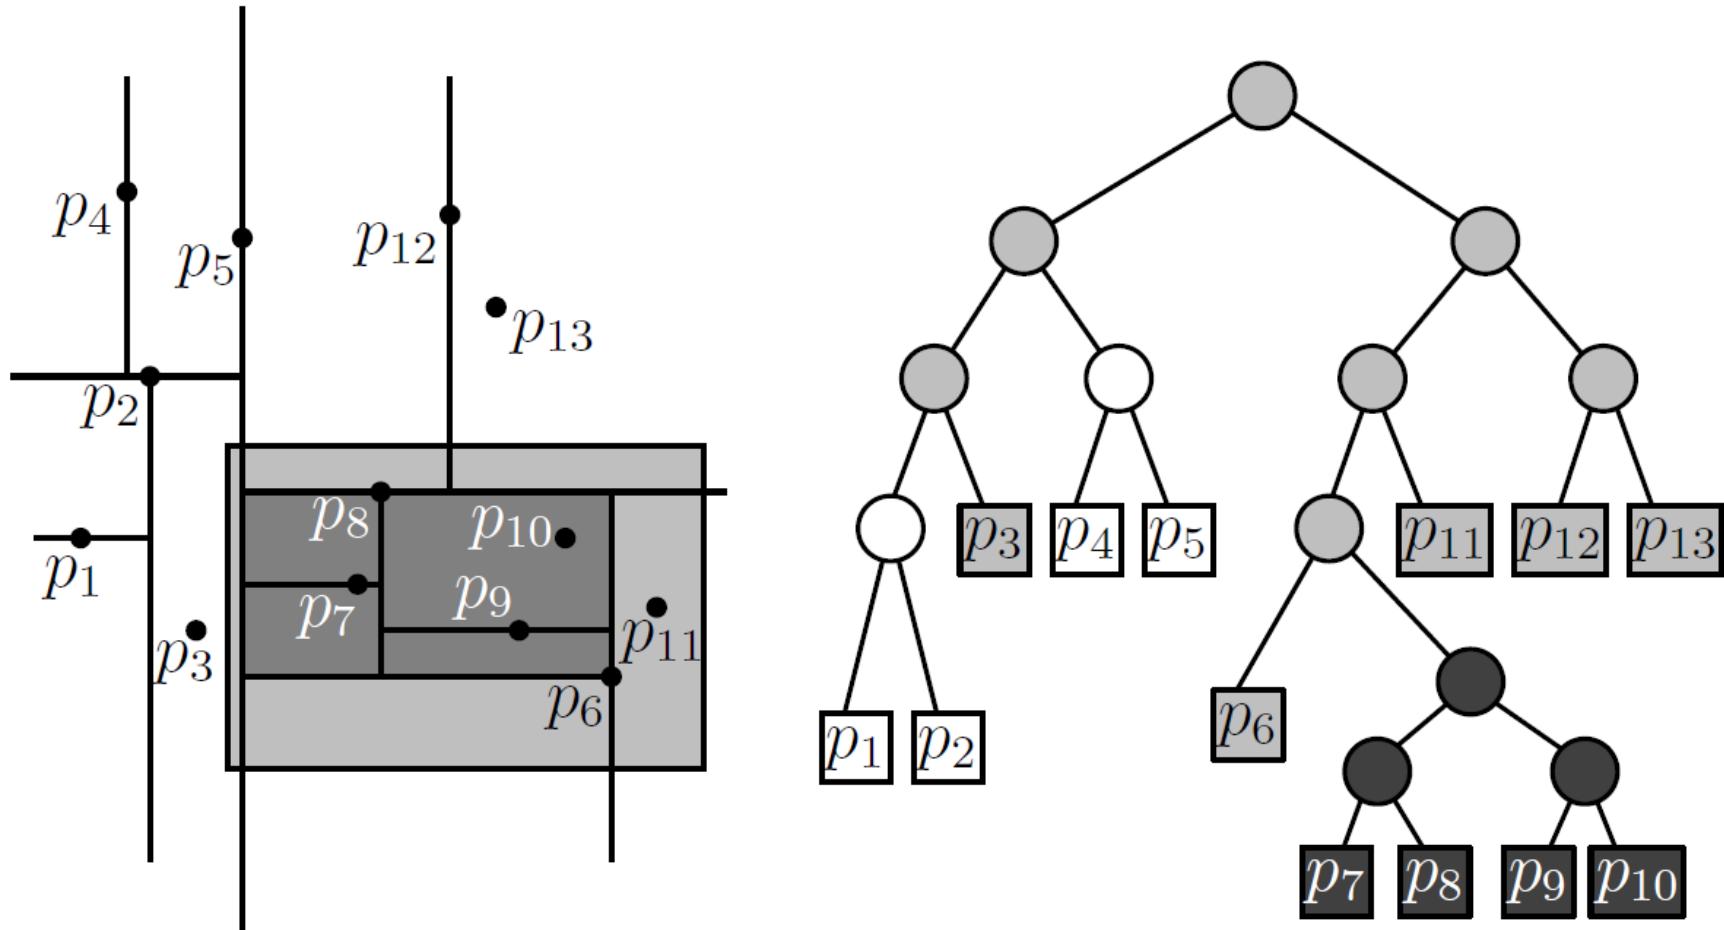
**Algorithm**  $\text{BUILDKDTREE}(P, \text{depth})$

1. **if**  $P$  contains only one point
2. **then return** a leaf storing this point
3. **else if**  $\text{depth}$  is even
4.     **then** Split  $P$  with a vertical line  $\ell$  through the median  $x$ -coordinate into  $P_1$  (left of or on  $\ell$ ) and  $P_2$  (right of  $\ell$ )
5.     **else** Split  $P$  with a horizontal line  $\ell$  through the median  $y$ -coordinate into  $P_1$  (below or on  $\ell$ ) and  $P_2$  (above  $\ell$ )
6.      $v_{\text{left}} \leftarrow \text{BUILDKDTREE}(P_1, \text{depth} + 1)$
7.      $v_{\text{right}} \leftarrow \text{BUILDKDTREE}(P_2, \text{depth} + 1)$
8.     Create a node  $v$  storing  $\ell$ , make  $v_{\text{left}}$  the left child of  $v$ , and make  $v_{\text{right}}$  the right child of  $v$ .
9. **return**  $v$

# Kd-Tree Region of Nodes



# Kd-Tree Querying



# Kd-Tree Querying

**Algorithm** SEARCHKDTREE( $v, R$ )

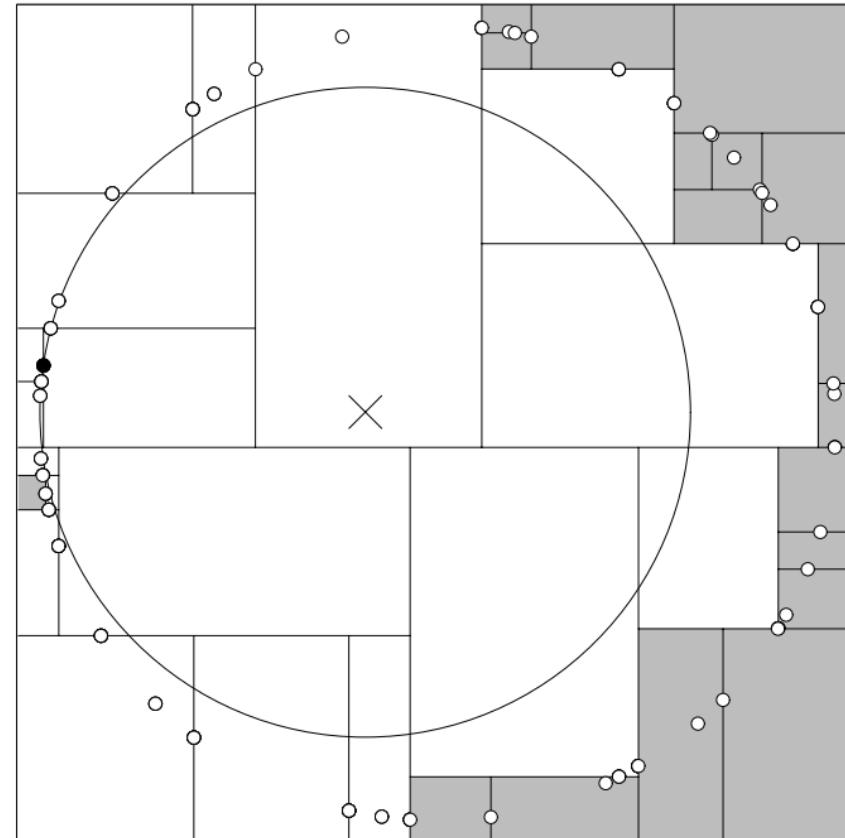
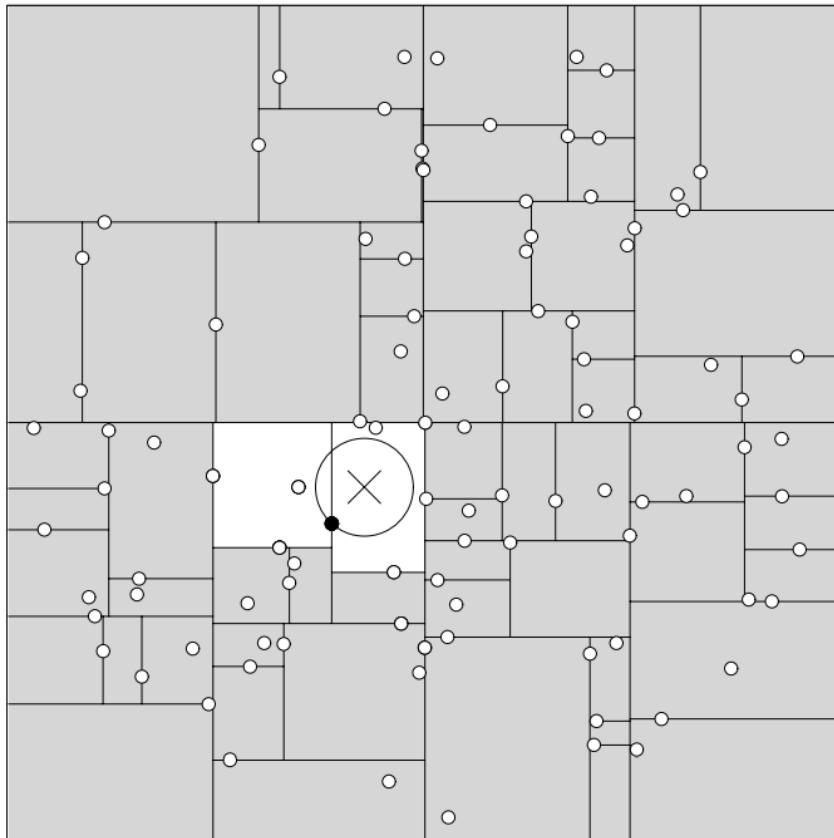
*Input.* The root of (a subtree of) a kd-tree, and a range  $R$

*Output.* All points at leaves below  $v$  that lie in the range.

1. **if**  $v$  is a leaf
2.     **then** Report the point stored at  $v$  if it lies in  $R$
3.     **else** **if**  $region(lc(v))$  is fully contained in  $R$   
4.         **then** REPORTSUBTREE( $lc(v)$ )
5.         **else** **if**  $region(lc(v))$  intersects  $R$   
6.             **then** SEARCHKDTREE( $lc(v), R$ )
7.         **if**  $region(rc(v))$  is fully contained in  $R$   
8.             **then** REPORTSUBTREE( $rc(v)$ )
9.         **else** **if**  $region(rc(v))$  intersects  $R$   
10.             **then** SEARCHKDTREE( $rc(v), R$ )

# Nearest Neighbor Search

- Effects of input distribution on search



# Kd-Trees in Higher Dimensions

Theorem: A set of  $n$  points in  $d$ -space  
can be preprocessed in  $O(n \log n)$  time into a data  
structure of  $O(n)$  size  
so that any  $d$ -dimensional range query can be  
answered in  $O(n^{1-1/d} + k)$  time, where  $k$  is the  
number of answers reported

# FLANN

- Adaptively selects between
  - Randomized kd-trees
  - Hierarchical k-means trees
- Criteria
  - Dimensionality
  - Relative weights on tree construction vs. search optimization

# Normal Estimation in Point Clouds

- Important task in itself
- Useful for encoding data in representations enabling classification, matching and indexing
- Assumption: each point lies on locally linear (planar) patch along with some (many) of its neighbors

# Normal Estimation in Point Clouds

- Find reference point's nearest neighbors
  - Forming triplets (ref. point + 2 NNs) increases complexity with dubious benefits
  - Form pairs and accumulate partial information, since two points do not define a surface
- Accumulate information for reference point's tangent plane
  - Vector connecting ref. point and each NN belongs to tangent plane

# Normal Estimation in Point Clouds

- Form scatter (covariance) matrix

$$\mathbf{M} = \frac{1}{k} \sum_{i=0}^k (\mathbf{p}_i - \hat{\mathbf{p}})(\mathbf{p}_i - \hat{\mathbf{p}})^T, \quad \hat{\mathbf{p}} = \frac{1}{k} \sum_{i=0}^k \mathbf{p}_i$$

- Normal: eigenvector with smallest eigenvalue (0 eigenvalue if perfect plane)

# Normal Estimation in Point Clouds

- Form scatter (covariance) matrix
- Normal: eigenvector with smallest eigenvalue (0 eigenvalue if perfect plane)
- Tombari et al. ECCV 2010:

$$M = \frac{1}{\sum_{i:d_i < R} (R - d_i)} \sum_{i:d_i < R} (R - d_i) (p_i - p)(p_i - p)^T$$

where,  $R$  is the radius of the ball around  $p$  and  $d_i$  is the distance from  $p$  to  $p_i$

# Normal Estimation in Point Clouds

- Tombari et al. resolve sign of eigenvectors ( $x, y, z$ ) by making them consistent with majority of neighborhood points

$$S_x^+ \doteq \{i : d_i \leq R \wedge (\mathbf{p}_i - \mathbf{p}) \cdot \mathbf{x}^+ \geq 0\}$$

$$S_x^- \doteq \{i : d_i \leq R \wedge (\mathbf{p}_i - \mathbf{p}) \cdot \mathbf{x}^- > 0\}$$

$$\mathbf{x} = \begin{cases} \mathbf{x}^+, & |S_x^+| \geq |S_x^-| \\ \mathbf{x}^-, & \text{otherwise} \end{cases}$$

- Where does the surface normal point for convex shapes?

# The Right Way

- Count nearest neighbors more, attenuate influence of remote points
  - At the very least, normalize vectors that contribute to covariance matrix
  - Use weight function on outer products that decreases with distance

# Normal Estimation

- My recommendation

$$M = \sum_{i:d_i < R} e^{-\frac{d_i^2}{2\sigma^2}} \frac{(p_i - p)(p_i - p)^T}{\|p_i - p\|^2}$$

with  $\sigma$  a parameter representing scale.

# Descriptors for 3D Point Clouds

# Invariant Descriptors

- Objective: represent point cloud (or surface) in a way that it can be compared and matched with other point clouds
- Trade-off between **Uniqueness** and **Repeatability**
- Invariance to:
  - Rigid transformation of the object
  - Viewpoint change of the scanner(s)
  - Sampling variations
  - Noise
- Local and global descriptors have been proposed
  - Focus on local here

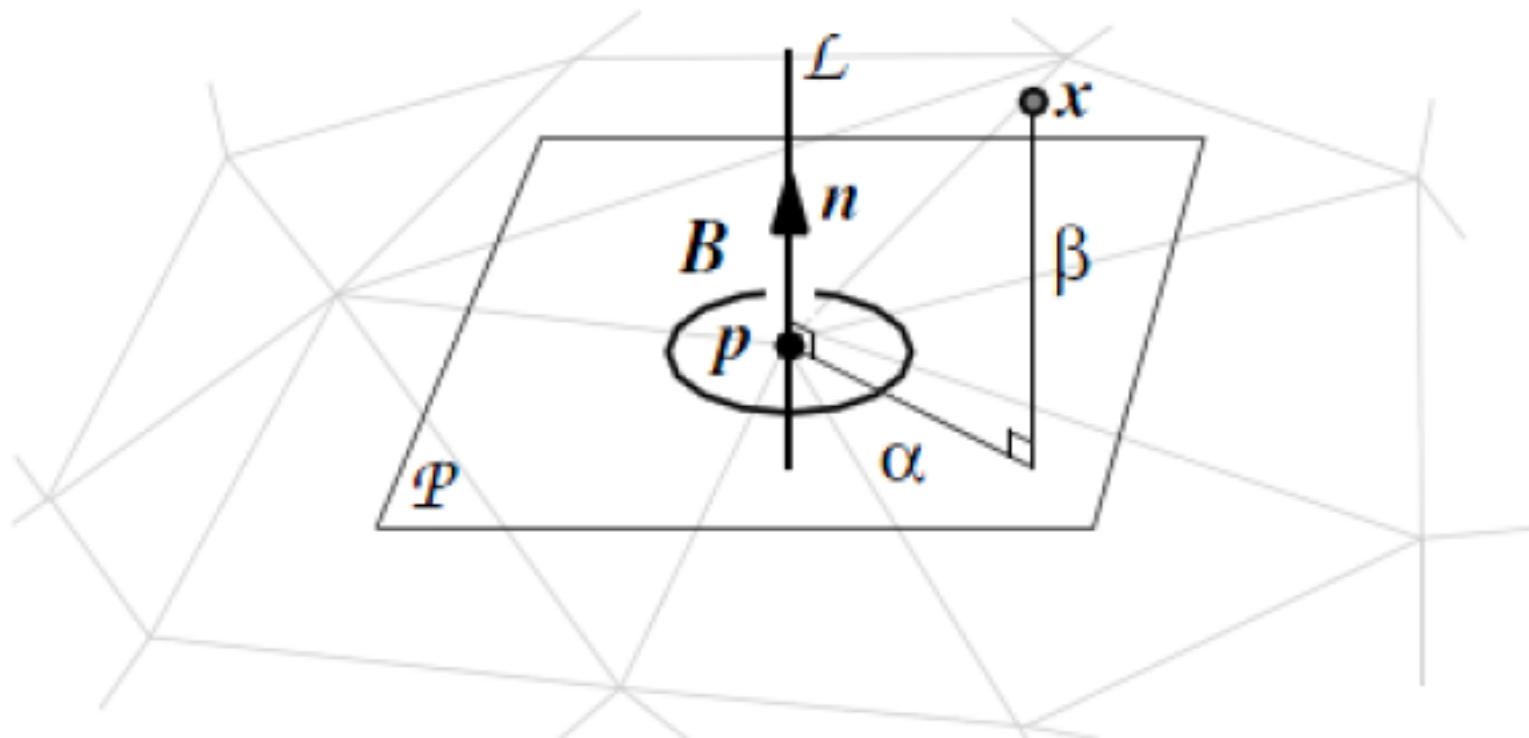
# Spin Images

- Johnson and Hebert, PAMI 1999
- Arguably most popular 3D shape descriptor, used in several recognition engines
- Fast to compute and very fast to compare

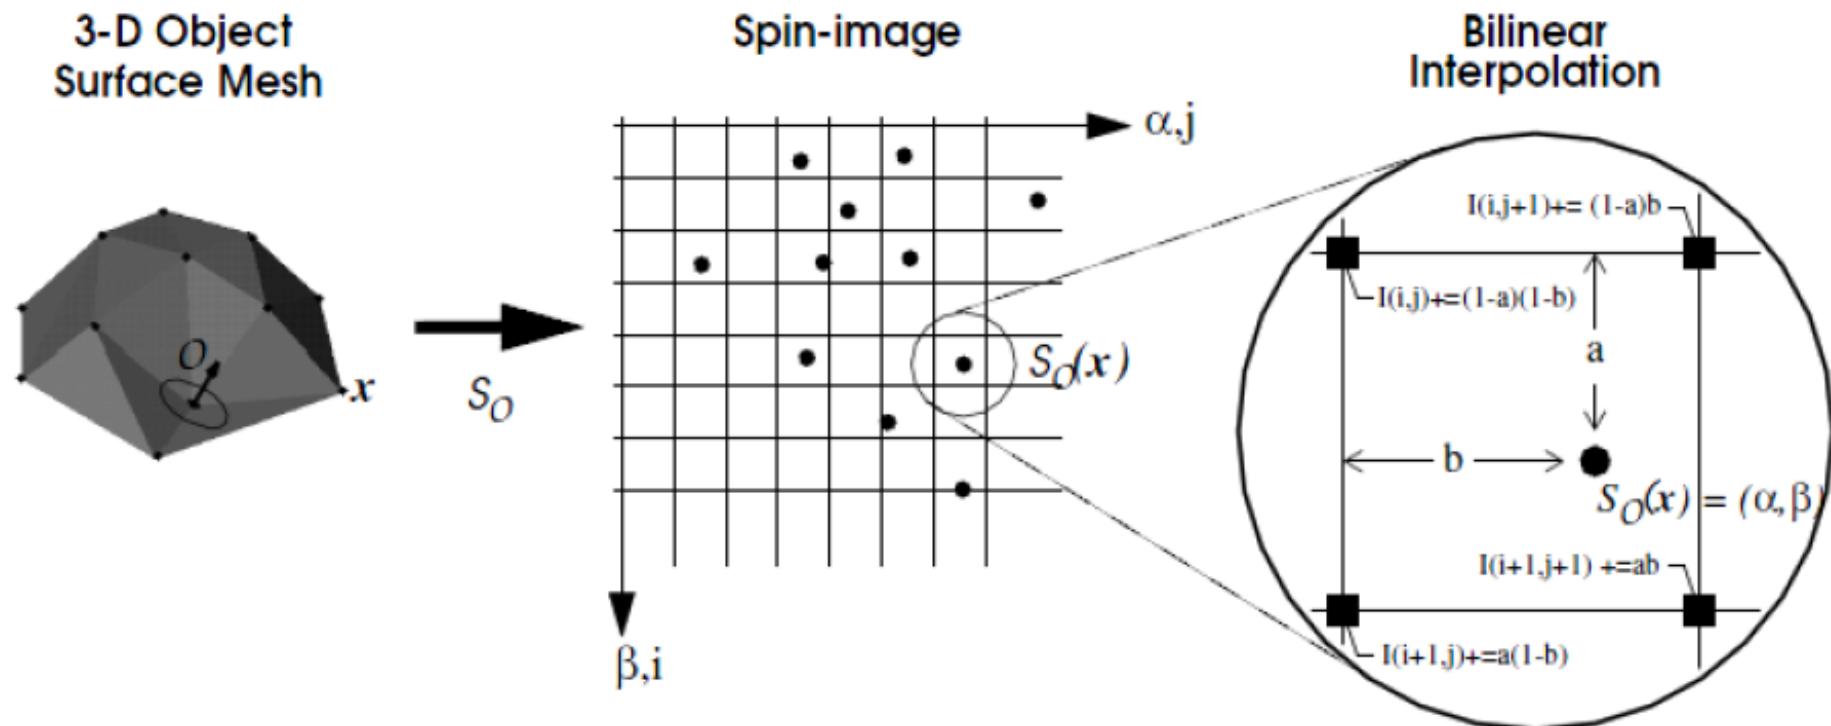
# Spin Images

- Computed in a cylindrical coordinate system defined by a reference point and its corresponding normal
- All points within this region are transformed by computing:
  - the distance from the reference normal ray  $\alpha$
  - the height above the reference normal plane  $\beta$
- A 2D histogram of  $\alpha$  and  $\beta$  is used as the descriptor
- Due to integration around the normal of the reference point, spin images are invariant to rotations about the normal

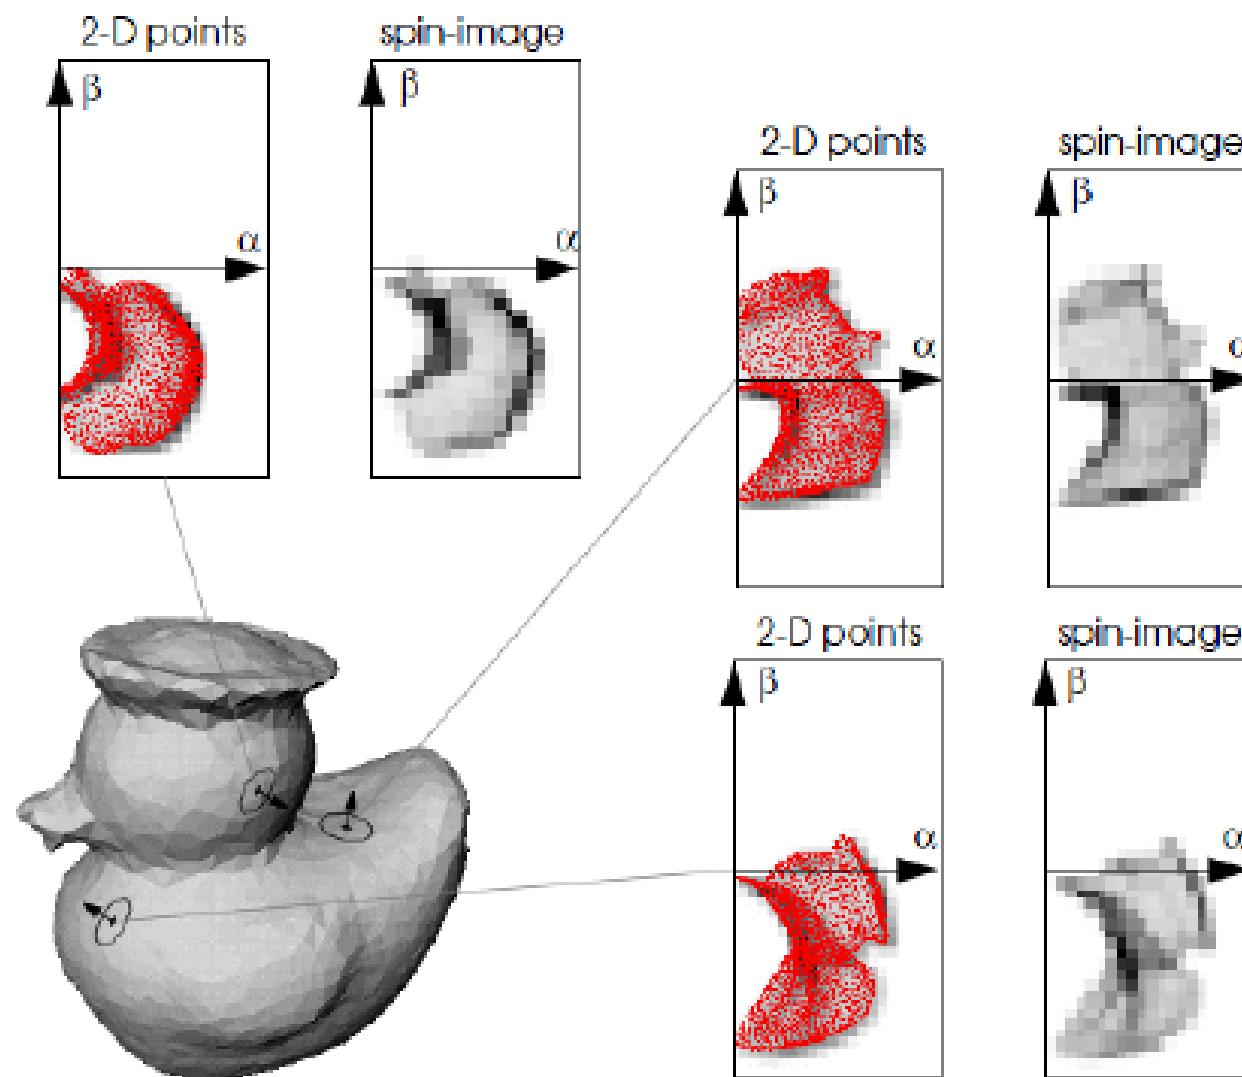
# Spin Images



# Spin Images

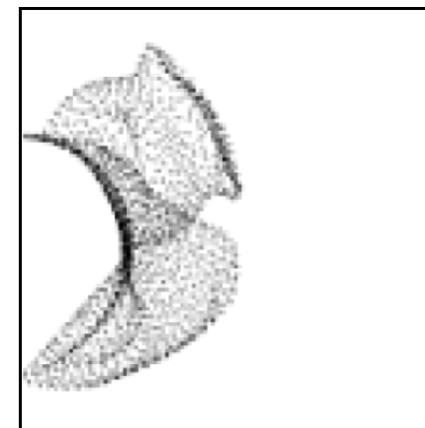
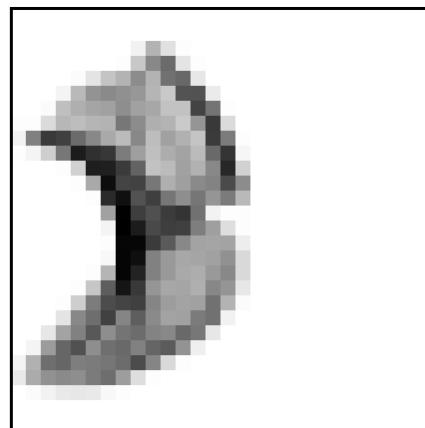
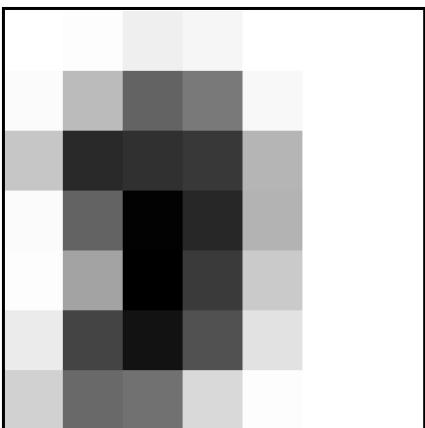


# Spin Images



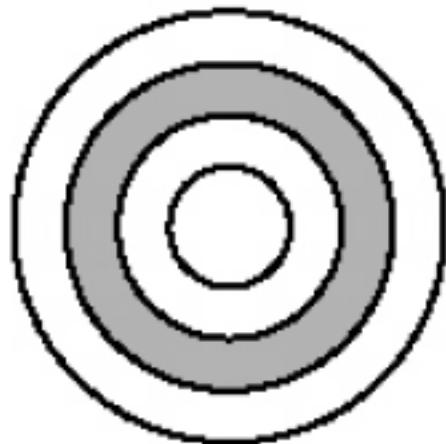
# Effects of Parameters

- Parameters of a spin image
  - Number of bins horizontally and vertically
    - Make vertical number of bins odd
  - Bin size
  - Support angle (max angle between reference normal and neighboring point normal to be included)
    - Default 60 degrees

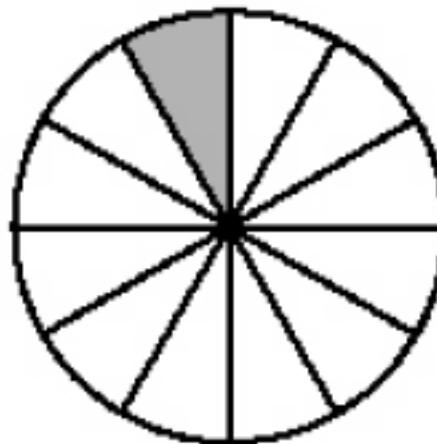


# 3D Shape Contexts

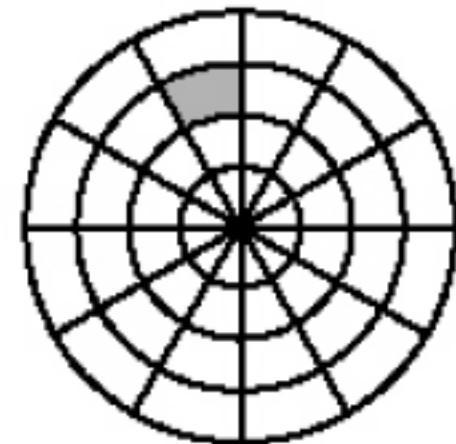
- Frome et al., ECCV 2004
- Histogram of neighboring points in sphere



4 shell bins



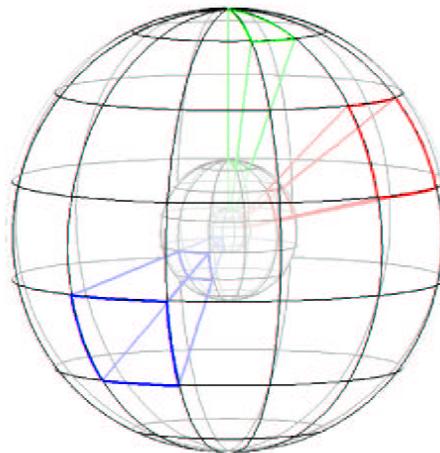
12 sector bins



48 combined bins

# 3D Shape Contexts

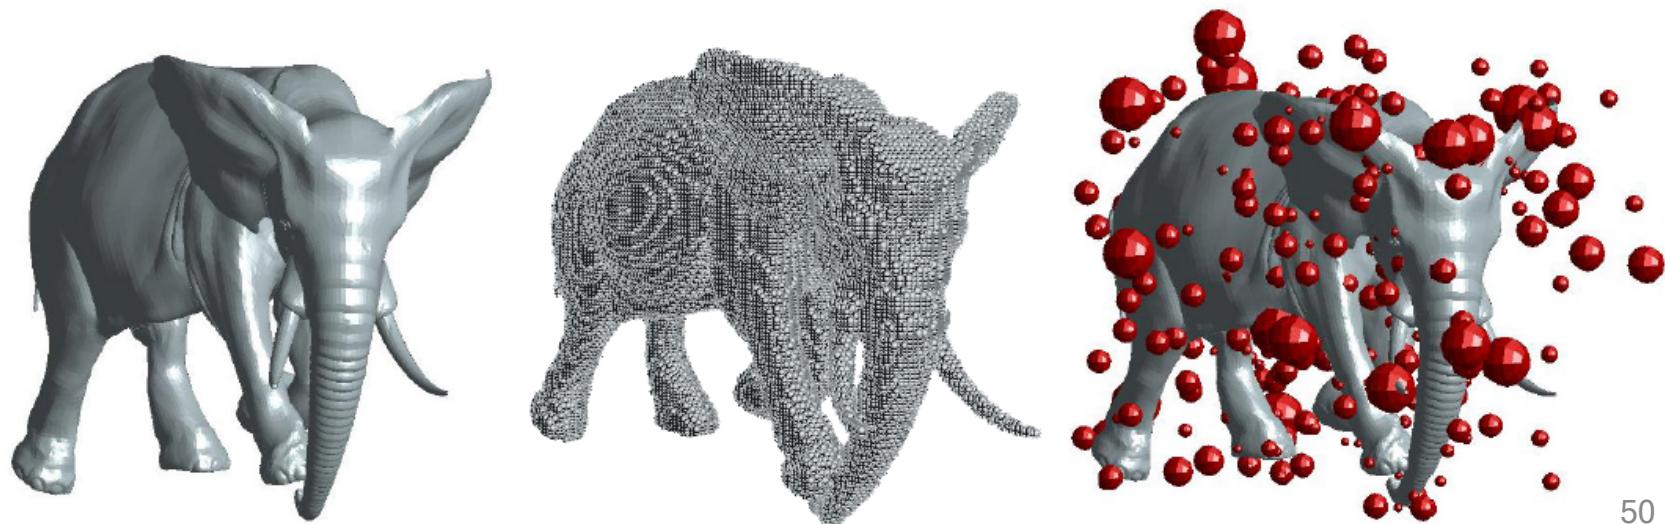
- Frome et al., ECCV 2004
- Histogram of neighboring points in sphere



- Challenge: matching requires rotations

# 3D SURF

- Knopp et al., ECCV 2010
- **Detector and descriptor**
- Convert surface into voxel representation
- Compute second-order derivatives at several scales (3 octaves)

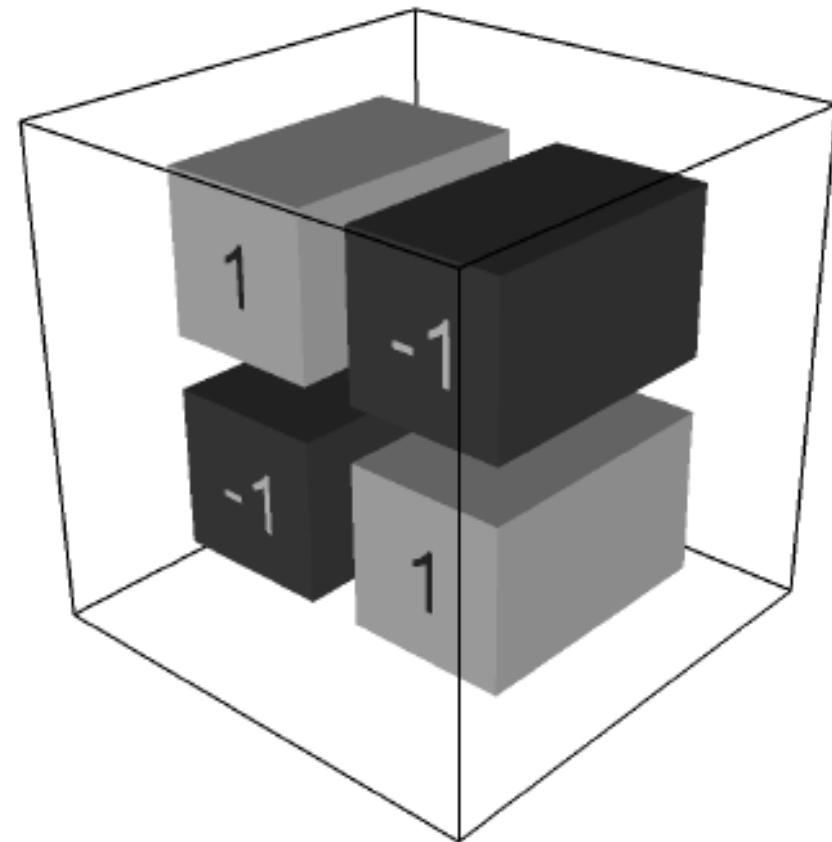
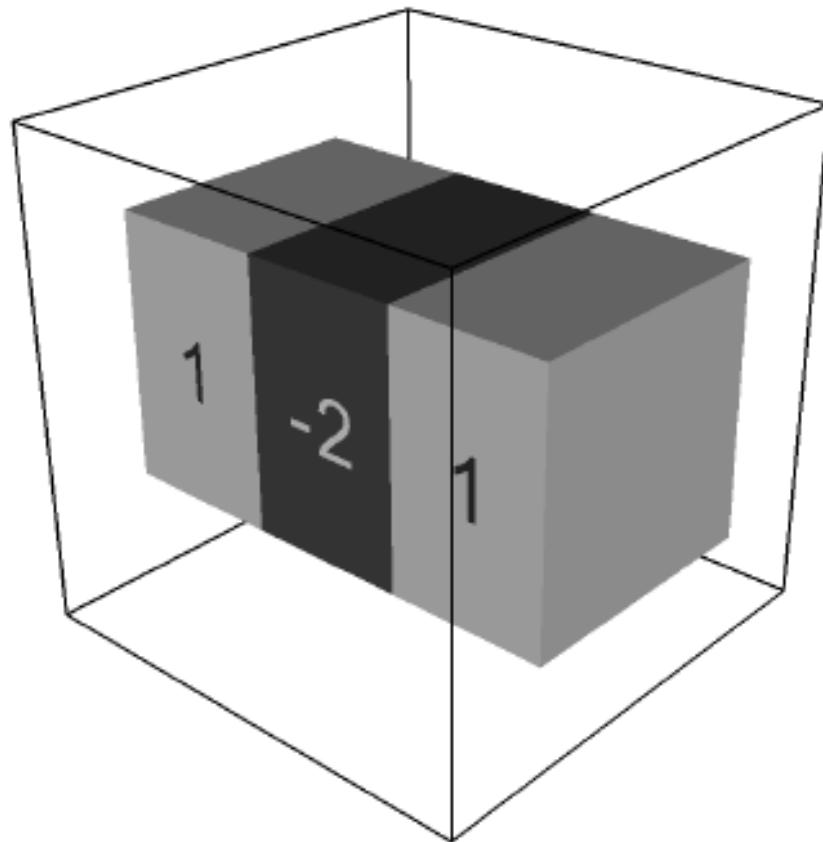


# 3D SURF

$$S(\mathbf{x}, \sigma) = |H(\mathbf{x}, \sigma)| = \left| \begin{pmatrix} L_{xx}(\mathbf{x}, \sigma) & L_{xy}(\mathbf{x}, \sigma) & L_{xz}(\mathbf{x}, \sigma) \\ L_{yx}(\mathbf{x}, \sigma) & L_{yy}(\mathbf{x}, \sigma) & L_{yz}(\mathbf{x}, \sigma) \\ L_{zx}(\mathbf{x}, \sigma) & L_{zy}(\mathbf{x}, \sigma) & L_{zz}(\mathbf{x}, \sigma) \end{pmatrix} \right|$$

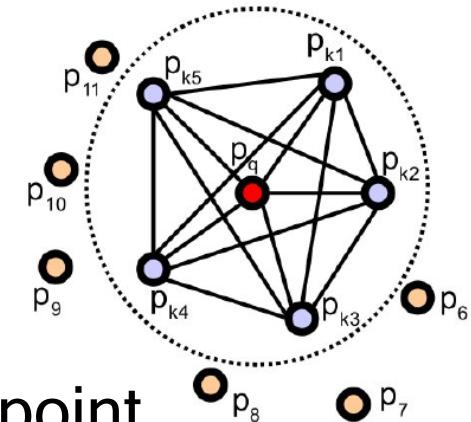
- Saliency function: absolute value of the determinant of the Hessian matrix at each point
- Select keypoints after non-max suppression
- Compute invariant local coordinate frame
- Descriptor:
  - $N \times N \times N$  grid around the feature
  - At each grid cell, store a 6-dimensional description vector of Haar wavelet responses
  - Default  $N=3$

# 3D Haar Filters



# Fast Point Feature Histograms

- Rusu et al., ICRA 2009 (other variations exist)
- PFH: Given points and normal:
  - Find all pairs of neighbors of reference point and define local frame
  - $u = n_i$
  - $v = (p_j - p_i) \times u$
  - $w = u \times v$
  - Compute properties of frame



$$\alpha = v \cdot n_j$$

$$\phi = (u \cdot (p_j - p_i)) / \|p_j - p_i\|$$

$$\theta = \arctan(w \cdot n_j, u \cdot n_j)$$

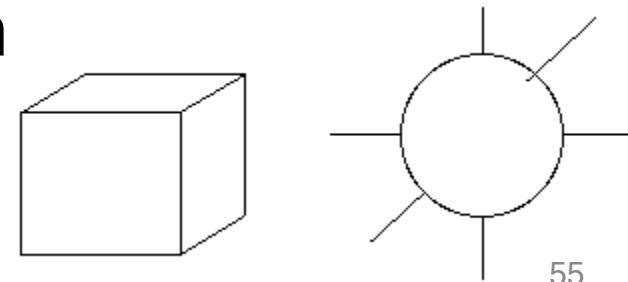
# Fast Point Feature Histograms

- PFH (cont)
  - Perform persistence analysis to determine which features are salient at a given scale
- FPFH:
  - Do not compute over all pairs of neighbors, but only between reference point and its  $k$  nearest neighbors
  - Then, blend Simplified Point Feature Histograms (SPFH) with weights inversely proportional to distances between points (typically 5 bins per dimension, 125-D descriptor)
  - More optimizations in paper

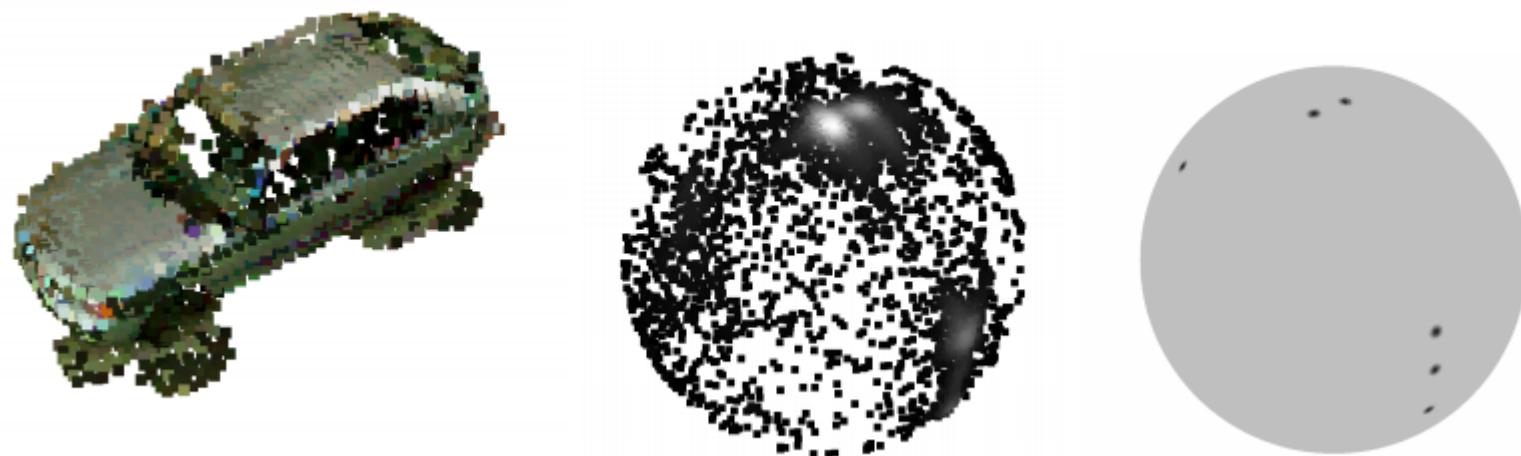
$$FPFH(p) = SPF(p) + \frac{1}{k} \sum_{i=1}^k \frac{1}{\omega_k} \cdot SPF(p_k)$$

# Extended Gaussian Images

- Horn, Proc. of the IEEE 1984
- Represent shape by mapping the normal of each point on unit sphere
  - Surface normal has two degrees of freedom
- Convex shapes can be uniquely reconstructed given EGI
  - Non-convex shapes can be described by EGI with some loss of information



# Extended Gaussian Images

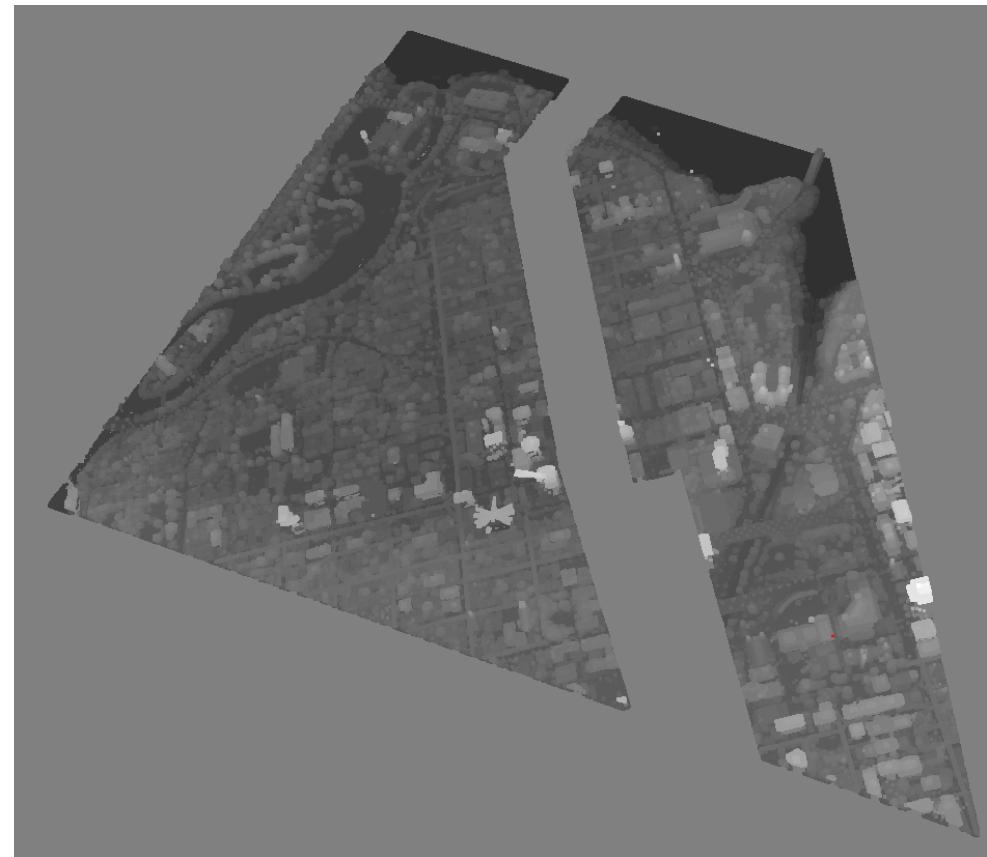


- Above: input point cloud, EGI, constellation EGI
- Matching requires alignment of two spherical histograms
  - Unpleasant, at best

# Analysis of Large-Scale 3D Point Clouds

# Segmentation of Large-scale 3D Datasets

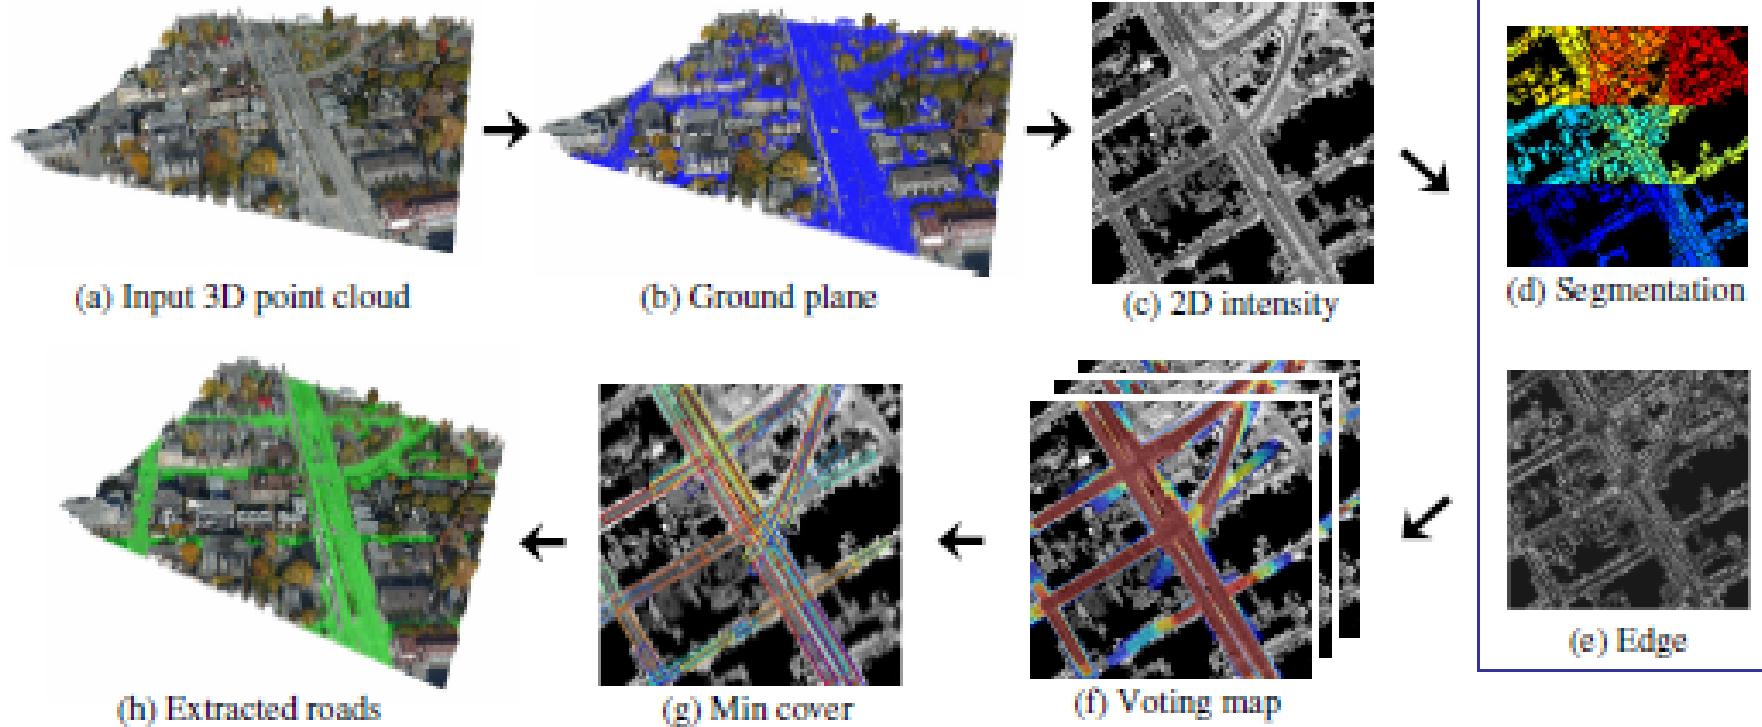
- Input: colored point clouds collected by terrestrial and airborne LIDAR sensors
- Goal: detect and recognize more than 100 objects classes
  - Stadiums and power plants
  - Mailboxes and parking meters
  - Powerlines



# A Minimum Cover Approach for Extracting the Road Network from Airborne LIDAR Data

- Combine local edge and region information to estimate road likelihood
- Pose road extraction as minimum cover problem
  - Explain likelihood maps by rectangular road segments with strong preference for elongated segments

# Overview of the Algorithm



# Hypothesis Generation

- Sample hypotheses from region boundaries returned by NCut segmentation on intensity image
  - Captures low-contrast boundaries
  - Test multiple values for width (10-25m) and length (40-300m)
- Output likelihood maps  $L$  for each width by combining boundary and interior feature strengths
  - Each likelihood map covers all orientations resulting in good performance at intersections

$$\mathcal{L}_k(x, y) = \sum_{H_i \in \mathcal{H}_k^{valid}} \beta_{bd} S_i^{bd}(x, y) + \beta_{int} S_i^{int}(x, y)$$

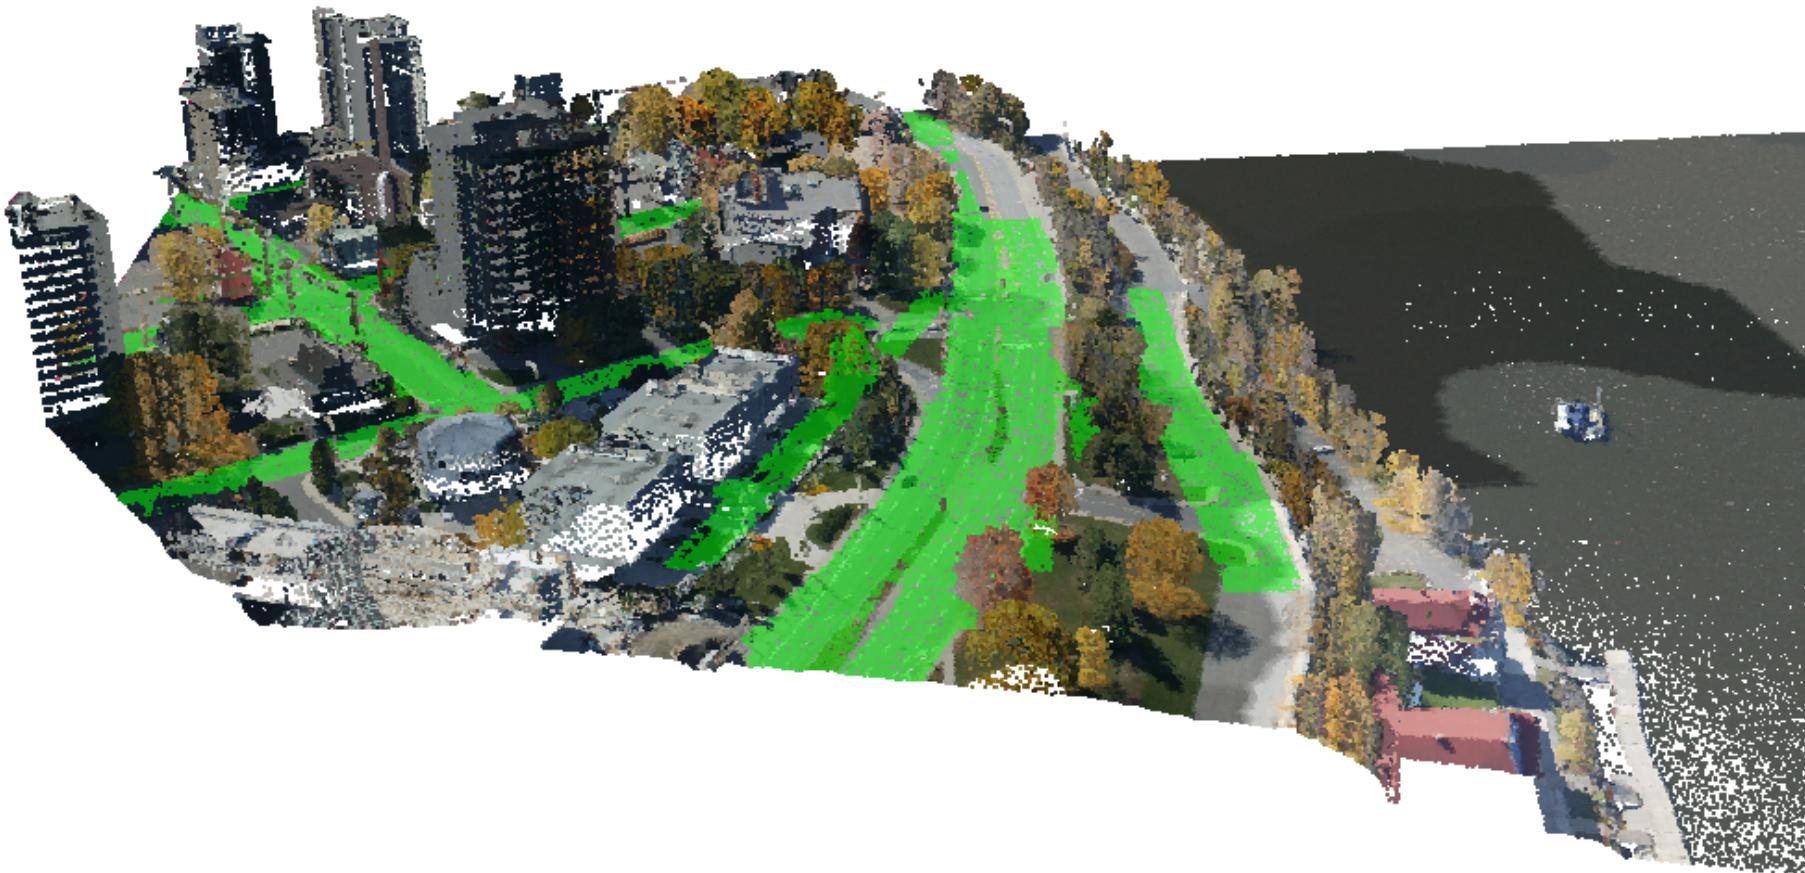
# Detection as Minimum Cover

- Explain likelihood maps by sparsest set of road segments
  - Penalize uncovered parts according to road likelihood
  - Penalize covered parts according to bg likelihood
  - Penalize for additional components
- NP-hard in general, but greedy approximation with theoretical guarantees is effective [Felzenszwalb and McAllester, 2006]

$$Cost(R_i) = \sum_{s_j \in R_i} C(s_j) + A$$

$$Cost_{cover}(\mathcal{R}_S) = \sum_{R_i \in \mathcal{R}_S} Cost(R_i) + \sum_{s_j \notin \mathcal{R}_S} \frac{\mathcal{L}(s_j)}{\mathcal{L}_{max}}$$

# Road Detection



# Road Detection



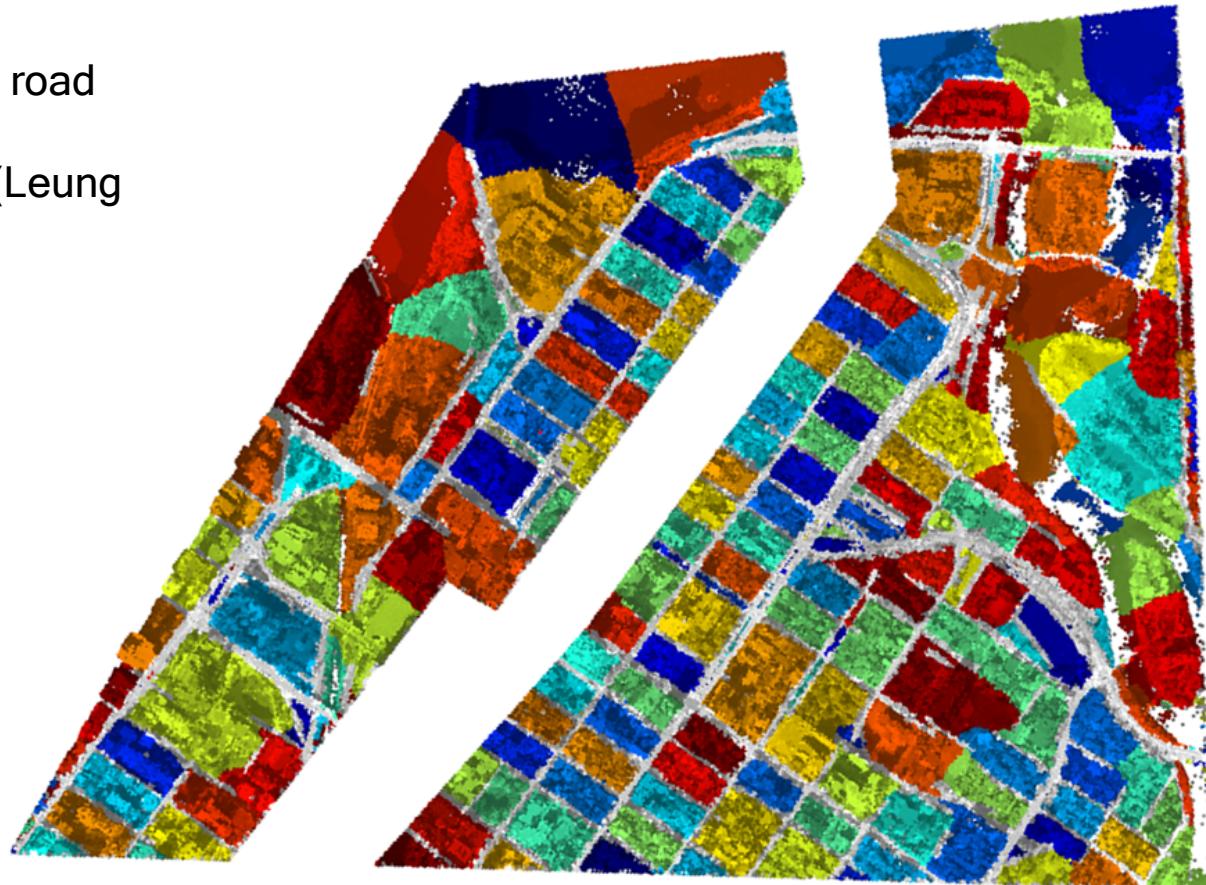
# Road Detection



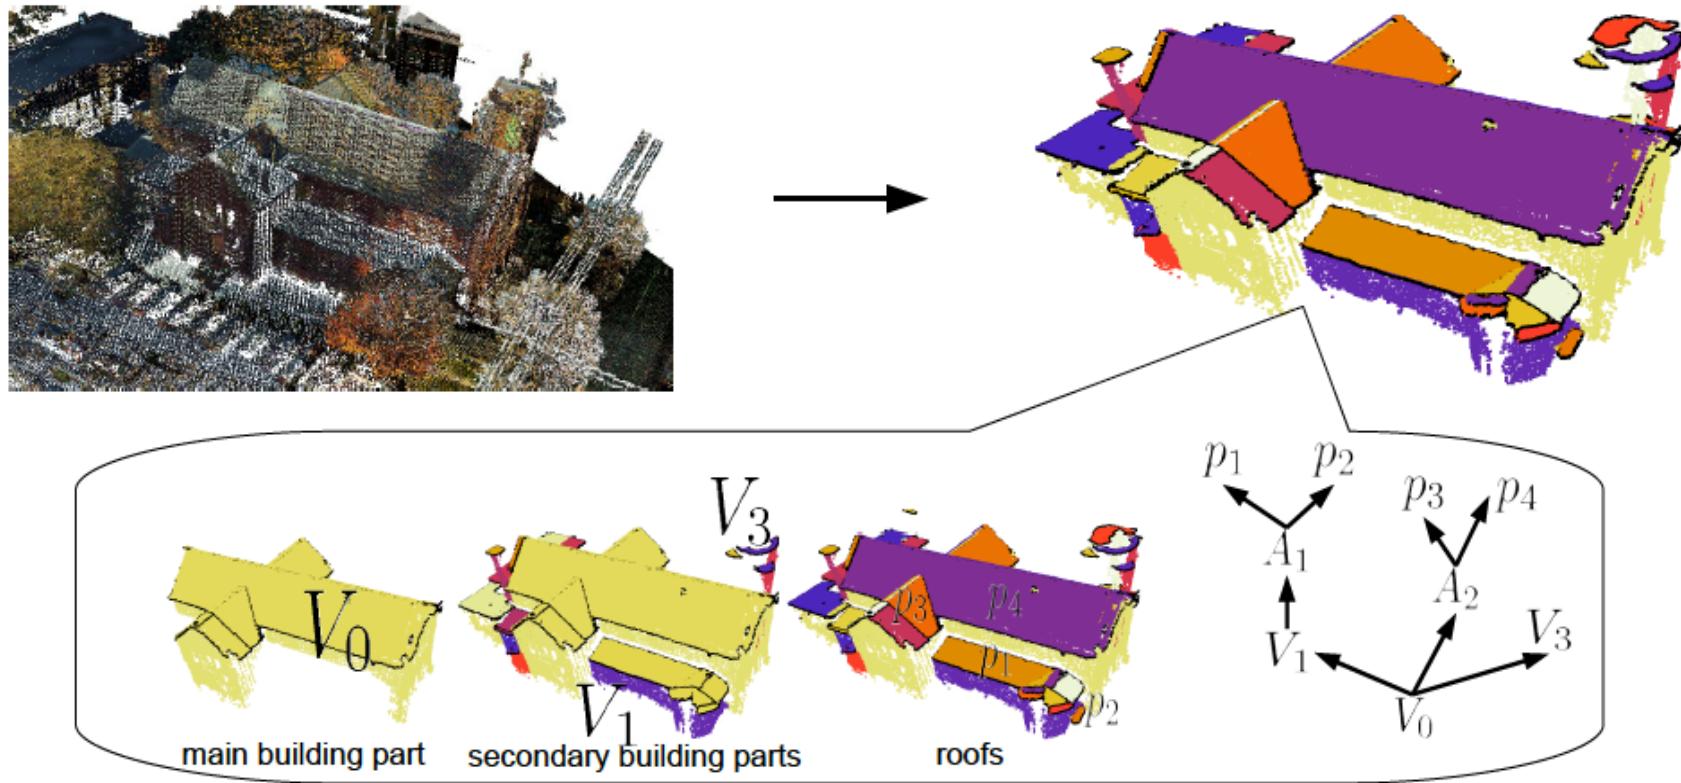
65

# City Blocks

Segment city blocks from road likelihood map using the intervening contour idea (Leung and Malik, 1998)



# Building Detection and Parsing



- Detection of buildings from unorganized range data
- Parsing of the buildings into an hierarchical, semantic representation

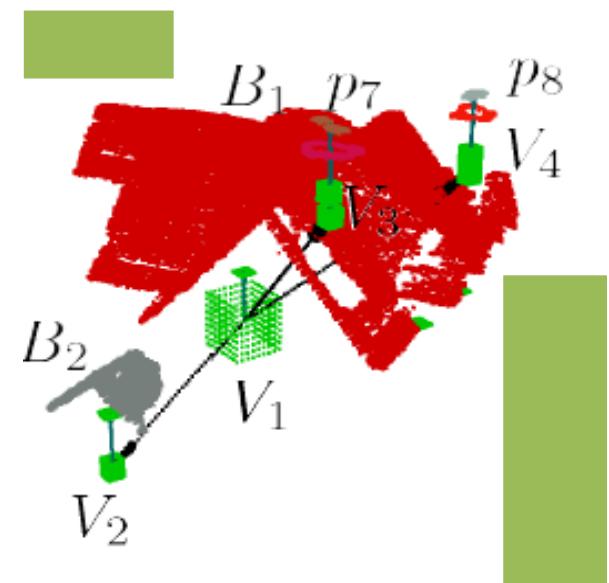
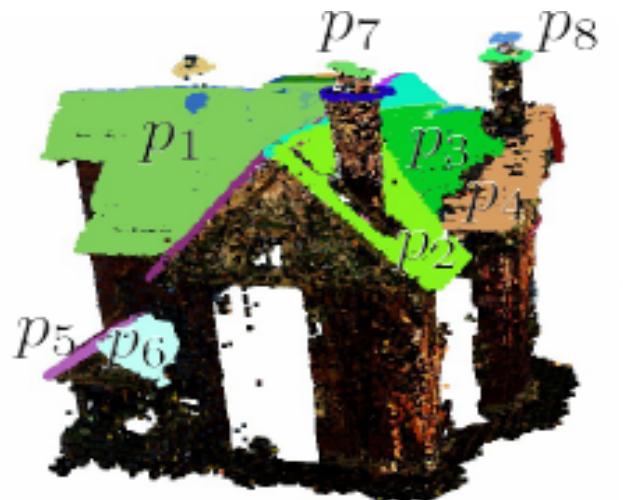
[Toshev, Mordohai and Taskar, CVPR 2010]

# Approach

- Use a generic grammar based on simple geometric rules
- Apply dependency parsing for efficient inference
- Define parsing and detection in a single framework
- Primitives:
  - Planes and planar patches - explicitly detectable from point clouds
  - Volumes - represent building parts and are enclosed by planar patches

# Grammar

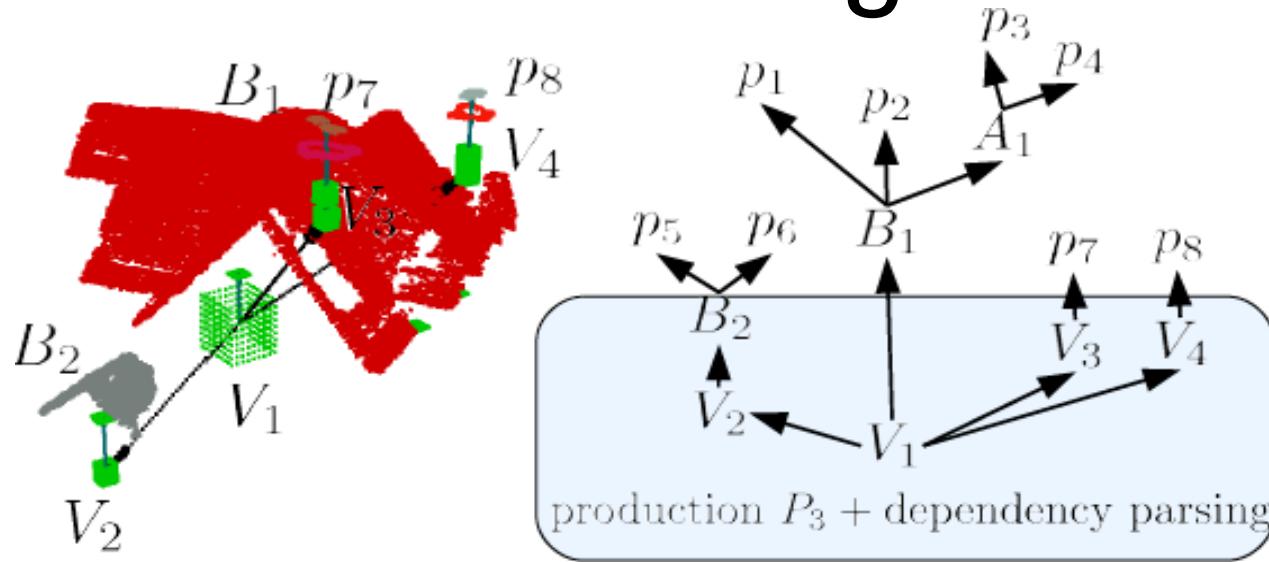
- Terminals: planar patches extracted from point cloud.
- Non terminals:
  - Roof components
  - Volumes enclosed by the roofs
  - Supernodes : a global “building” and a global “non-building” node used for detection



# Classification

- Classification of a volume using a linear SVM
- Features are extracted from planar patches enclosing the volume:
  - Elevation
  - Distance to the nearest ground point
  - Convexity of the upper volume surface
  - Scatter of point cloud in the volume
  - Area and aspect ratio
  - Degree of enclosure by empty space
  - Fitting error

# Parsing



- Productions for parsing planar patches are deterministic
- Hierarchy among volumes is not deterministic - dependency parsing
- For set of planar patches, a sequence of productions generates a parse tree

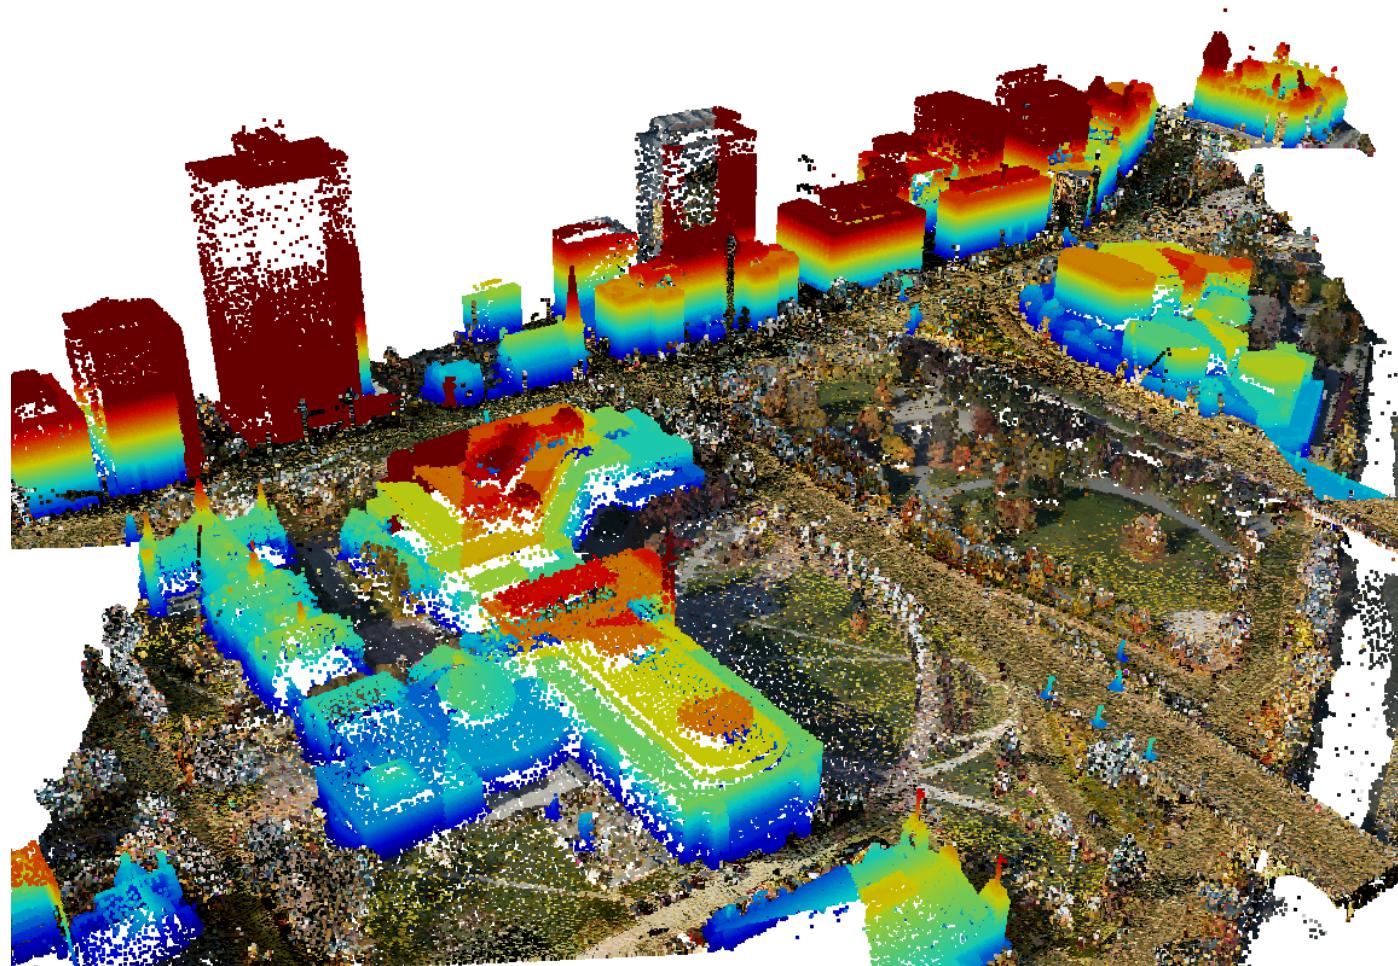
# Inference

- Construct a directed graph  $G$  consisting of the volumes
- Edge weights represent:
  - Hierarchy based on the area
  - How likely they belong to the same class
  - Whether they are children of the supernodes
- A parse tree  $T$  is a maximum spanning tree in  $G$
- Use Chi-Liu/Edmonds algorithm to compute MST

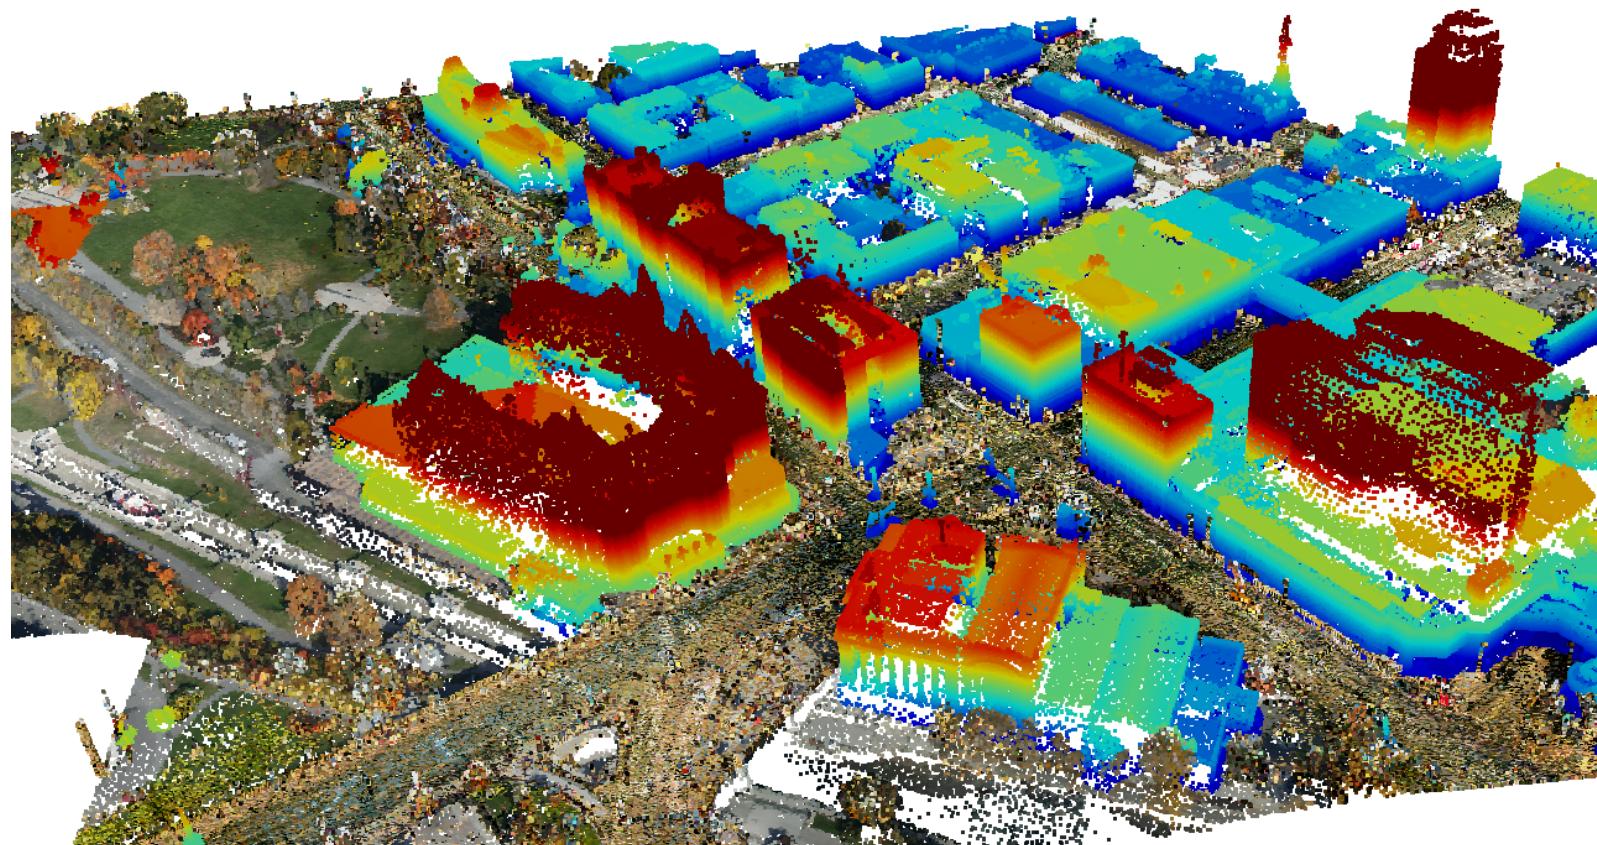
# Quantitative Results

- 9 blocks used for training (buildings and their parses are labeled)
- 78 blocks used for testing (only buildings are labeled)
- Detection results (accuracy of patch classification): 89.3%
- Parsing accuracy (3-fold cross-validation on the 9 blocks): 76.2 %

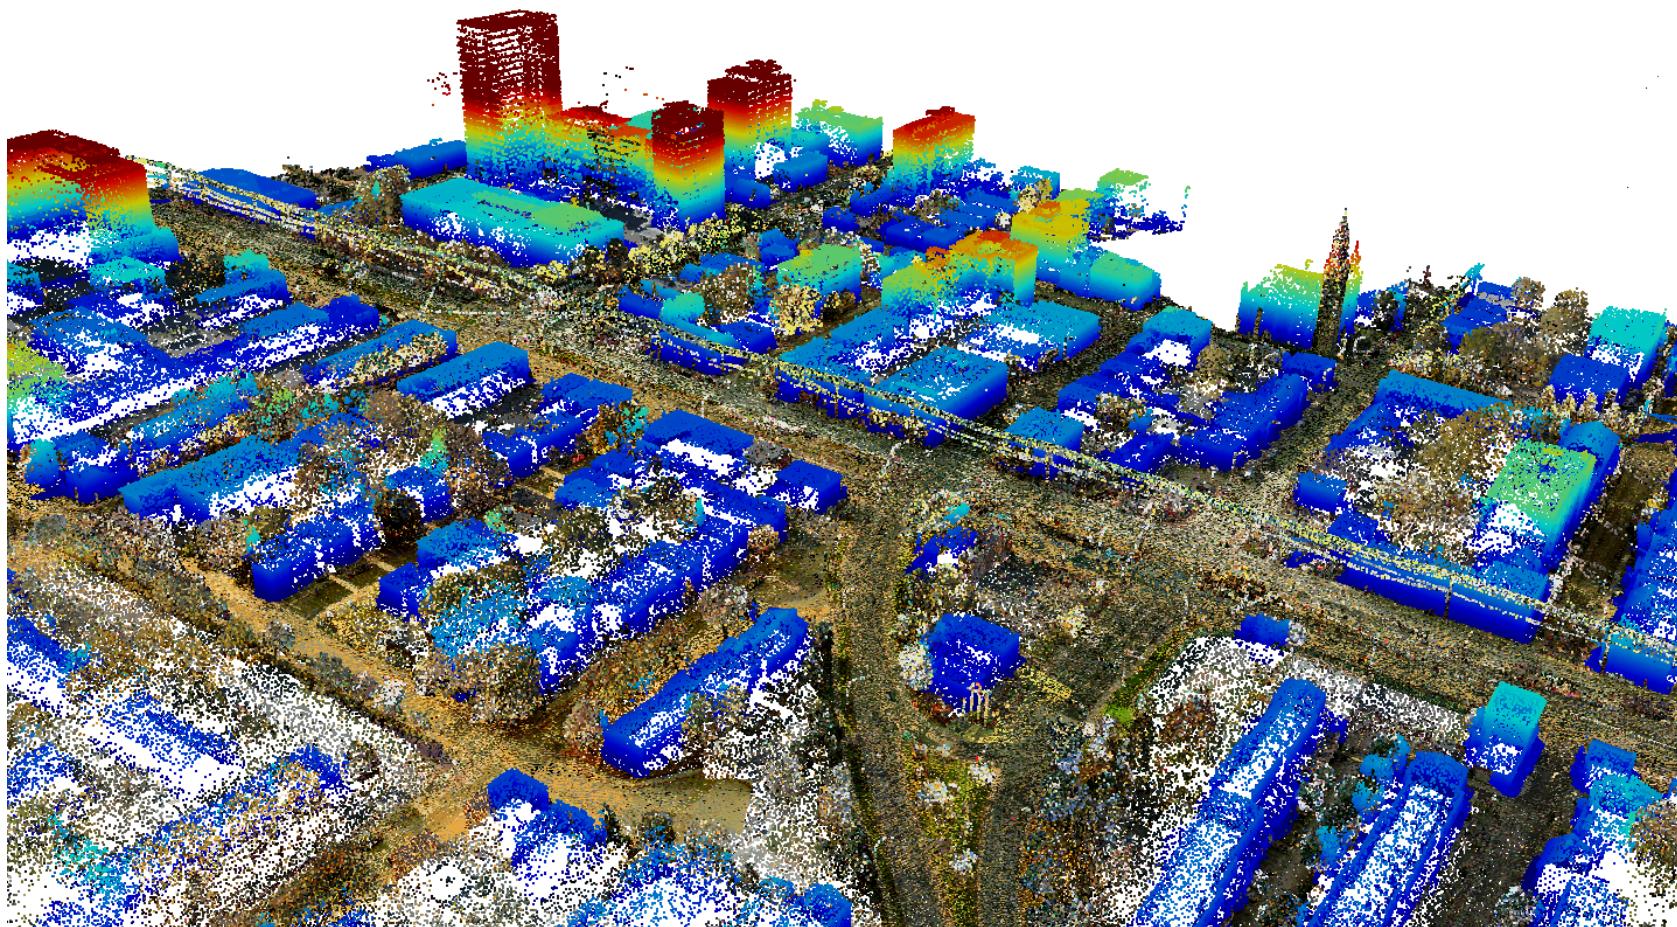
# Detection Results



# Detection Results

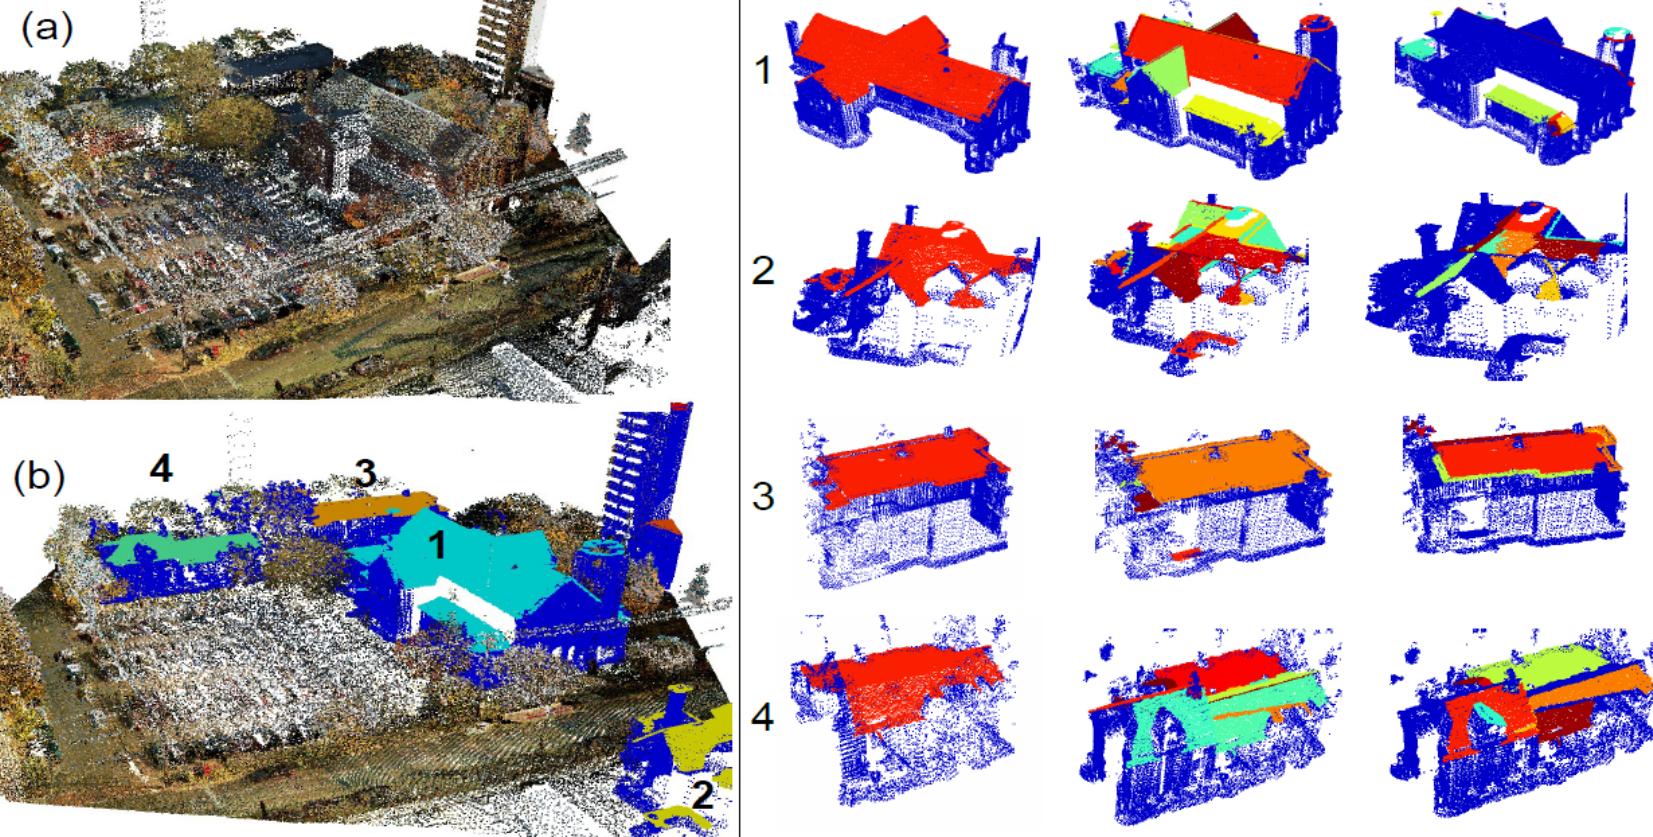


# Detection Results



# Parsing Results

original point cloud and detected buildings



# 3D Object Detection using Bottom-up and Top-down Descriptors

- Detect objects in large-scale 3D datasets
- Requirements:
  - Precision
  - Recall
  - Speed



[Patterson, Mordohai and Daniilidis, ECCV 2008]

78

# Shape Descriptors

- Different descriptors provide different trade-off between speed and accuracy
- More types of invariance => faster, less discriminative
  - E.g. spin images vs. 3D shape contexts
- Global descriptors are more accurate, but are sensitive to occlusion (and deformation)
  - Require only one comparison per target
  - Need segmentation hypotheses

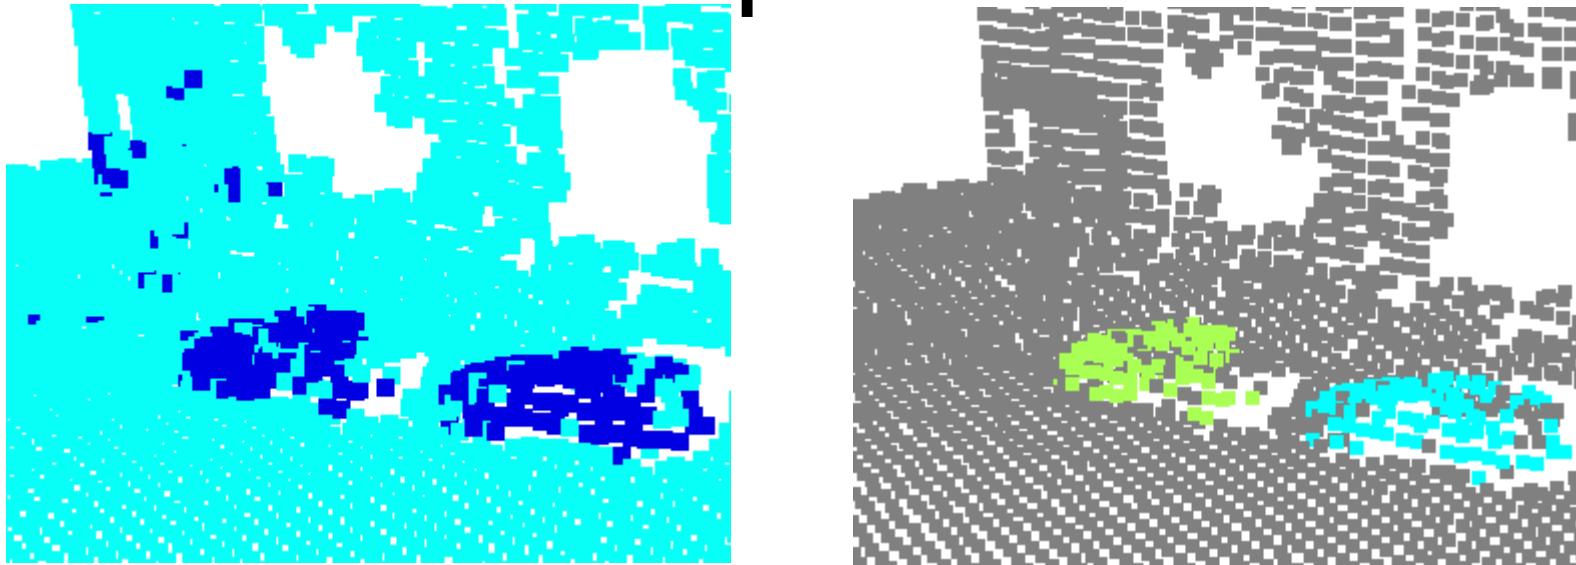
# Approach

- Spin images for fast bottom-up detection of regions of interest
  - Objective: maximize recall
- Extended Gaussian Images (EGIs) as global top-down descriptors to verify hypotheses
  - Objective: prune wrong hypotheses
  - Side product: best alignment with most similar model

# Experiments

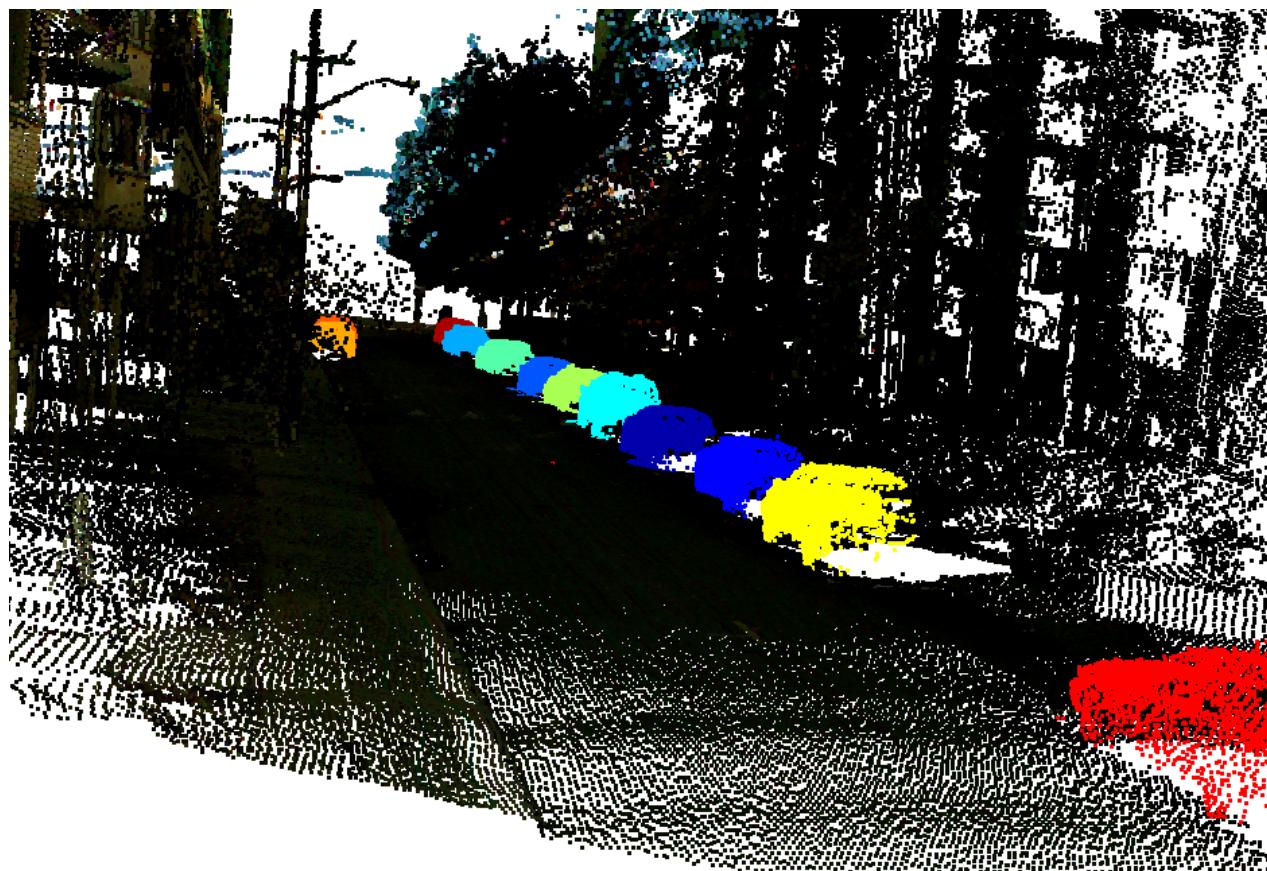
- Dataset: 220 million points collected from terrestrial sensors
- 2.2 million used as training set for spin images
  - 81,000 spin images in DB<sub>SI</sub>, 2600 positive
  - 17 cars in DB<sub>EGI</sub>

# Bottom-up Detection

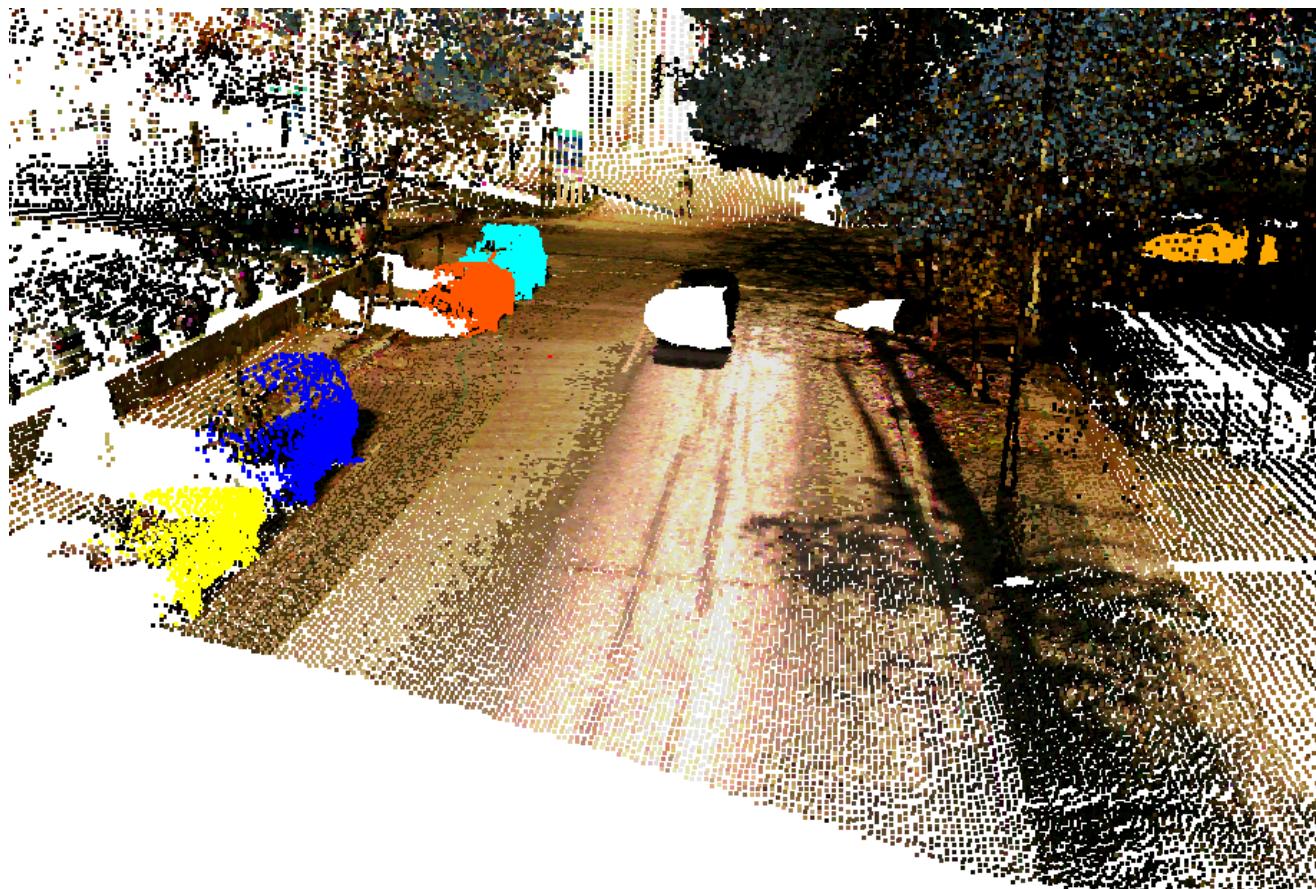


- Spin images of unknown scene classified as positive or negative
- Positive spin images clustered to form hypotheses
  - Minimum number required for hypothesis

# Results



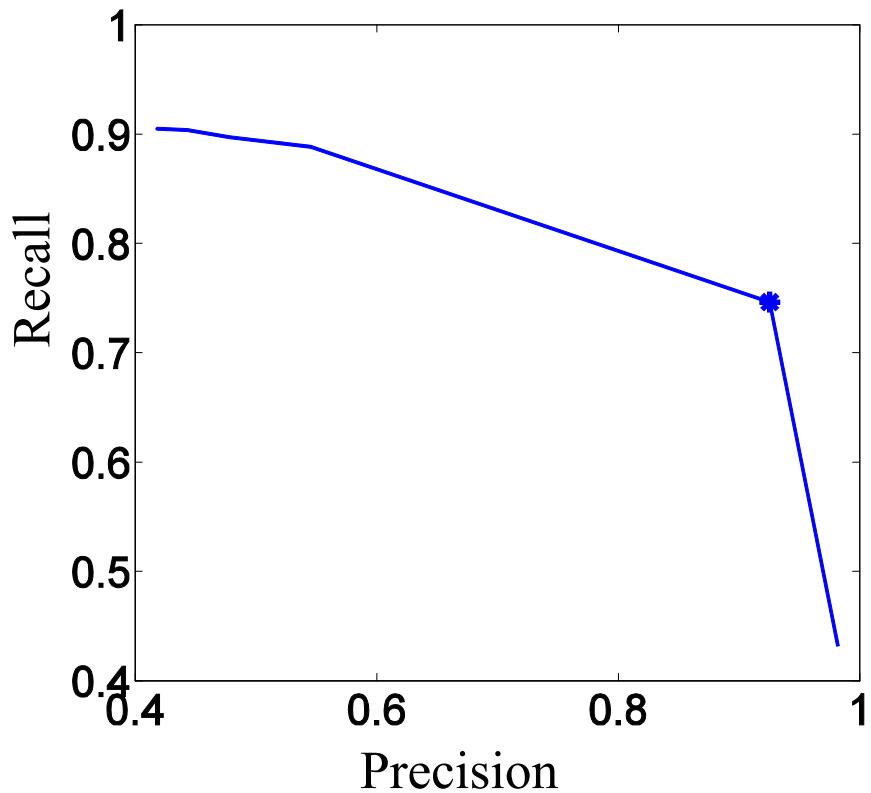
# Results



# Results



# Precision-Recall Curve



- Data: 1221 cars
- Bottom-up stage:
  - 2200 hypotheses,  
1100 correct
- Top-down stage \*:
  - 905 true positives
  - 74 false alarms
  - 316 missed detections

# Car Detection Video

